

The Incommensurability Principle in Biological Transport: Scale-Free Formulation, Topological Rigidity, and the Minimax Origin of Vascular Scaling

Riccardo Marchesi

University of Pavia

June 27, 2026

Abstract

Why does the mammalian vascular tree maintain a conserved branching exponent $\alpha^* \approx 2.72$ across a 10^7 -fold range in body mass, despite a fundamental shift in the underlying physics from viscous to wave-dominated transport? We demonstrate that this universality cannot emerge from local optimization under symmetric rules, as any junction-level coupling of incommensurable costs would require scale-dependent fine-tuning varying by $O(10^2\text{--}10^3)$ across the hierarchy—a biologically implausible constraint. Real vascular networks resolve this bound through structural heterogeneity. We show that vascular geometry emerges as a scale-free attractor of a *network-level minimax* principle.

By grounding the fitness penalty in ATP stoichiometry, we derive a scale-free cost functional and prove a Topological Rigidity theorem: the optimal branching exponent depends only on dimensionless structural parameters (G, N, p, α_w) and is independent of all pure metabolic quantities—blood oxygen cost, cardiac output, segment length, ATP stoichiometry. A self-consistency condition on the viscous–inertial energy partition at each bifurcation yields a dual-threshold framework with fluid threshold $Wo_c^{\text{fluid}} = \sqrt{3}$ and wave threshold $Wo_c^{\text{wave}} = 3/\sqrt{2}$ in mammalian vascular trees. The symmetric model derived standalone here yields $\alpha_{\text{model}}^* \approx 2.629$, in quantitative agreement with mammals near the allometric transition; scale-dependent morphometric heterogeneities (resolved by the coherent impedance solver of the companion work) shift large-mammal values toward 2.72. The framework explains the developmental stability of cardiovascular networks as a consequence of the architecture being decoupled from the biochemistry.

1 Introduction

The heart of a mouse (*Mus musculus*) beats over 600 times per minute, while the heart of a blue whale (*Balaenoptera musculus*) beats barely 20 times. These organisms differ by seven orders of magnitude in mass, and the fluid dynamics within their coronary arteries inhabit vastly different physical regimes: the whale’s transport is dominated by massive inertial waves, while the mouse’s is governed by the viscous linearity of Murray’s law. What sets these regimes apart is not the beat frequency but the volume of blood each heart drives through its arteries: cardiac output scales with body mass as $\dot{Q} \propto M^{3/4}$ [1] and stroke volume as $\propto M$, so a single whale contraction ejects a blood volume larger than the mouse’s by the same seven orders of magnitude. This volume is carried by vessels whose radius scales as $R \propto M^{3/8}$, and it is this radius—not the rapid murine heart rate—that dominates the Womersley number $Wo = R\sqrt{\omega/\nu}$ (with $\omega = 2\pi f$). Since $Wo \propto M^{3/8}\sqrt{M^{-1/4}} = M^{1/4}$ (derived in §6), it *increases* with body size

despite the whale’s slower rhythm. In isolation, each regime selects a different single-mechanism optimum of the unified network Lagrangian developed below: the wave-impedance attractor ($\alpha_w = 2.000$) for the inertia-dominated whale, and the viscous transport optimum ($\alpha_t \approx 2.920$, the network-level analogue of Murray’s law) for the mouse. Yet, their coronary branching exponents are nearly identical ($\alpha^* \approx 2.72$), and their vascular architectures exhibit a conserved geometric self-similarity that defies the massive shift in metabolic and hydrodynamic scales.

This structural universality poses a deep challenge to biological optimization theory. If vascular remodeling were governed by purely local rules—junction-level adaptations to metabolic demand and wave reflections—the universality of α^* would require an unphysical degree of coordination. As we prove in Section 2, any local coupling capable of preserving a fixed α^* across masses and generations requires a sensing mechanism that “knows” the global scale of the organism. Such a solution under symmetric local rules is not merely fine-tuned but requires scale-dependent coupling $\mu(g)$ varying by $O(10^2\text{--}10^3)$ across the hierarchy—a biologically implausible degree of fine-tuning. Real vascular networks resolve this constraint through structural heterogeneity (asymmetry, taper, generation-dependent gene expression).

In this work, we propose that the vascular tree avoids this informational cost by converging to a scale-free *network-level minimax* attractor. We demonstrate that the incommensurability—the dimensional incompatibility between metabolic costs (measured in Watts) and wave-reflection costs (dimensionless), analogous to geometric frustration in physical systems—forces the optimization to a higher topological level. By grounding the fitness landscape in the fundamental *ATP stoichiometry* of metabolism (Section 3), we derive a unique scale-invariant penalty functional that is immune to metabolic fluctuations.

Remark 1 (Paper Series Structure). This manuscript is the fourth in a series of companion works establishing the theoretical foundation, computational framework, and empirical validation of the incommensurability principle. For brevity and mathematical cohesion, we refer to: **Paper I** [2] for the static transport optimum and Murray’s law derivation; **Paper II** [3] for the coherent wave-reflection model and minimax formulation; **Paper III** [4] for the metabolic scaling theorem ($\beta_M = 3/4$) and allometric coupling. All theoretical results in this work are self-contained, with cross-references provided only for detailed derivations exceeding the scope of a single manuscript.

Critical Falsifiable Predictions

This framework makes three quantitative predictions testable with existing experimental techniques:

1. **Hummingbird Heart-Rate Override:** A 4 g hummingbird ($HR \approx 1000$ bpm) operates at $Wo \approx 1.76$ despite its small mass, yielding $\alpha^* \approx 2.6$ (transition regime, far from the Murray viscous limit $\alpha \approx 3.0$ expected from its mass). This prediction contradicts mass-only scaling models.
2. **Ontogenetic Phase Transition:** Mouse embryonic vasculature should exhibit $\alpha \approx 3.0$ at E8–E10 ($M < 0.1$ g, viscous regime), transitioning to $\alpha^* \approx 2.7$ by E15–E18 ($M > 1$ g, wave regime).
Falsification criterion: If $\alpha \approx 2.7$ from the earliest angiogenesis stages, the theory is falsified.
3. **Retinal Dual-Attractor Paradox:** 2D retinal vessels exhibit simultaneous divergence: diameter scaling $\alpha_{dia} \approx 2.0$ (2D wave-matching), bifurcation angles $\theta \approx 75.0^\circ$ (3D Murray equilibrium). This dual state is impossible under local optimization (Section 8).

The transition between viscous and wave regimes is governed by a fundamental topological

condition at the branching junctions. We establish the following criterion through a three-level derivation combining deductive geometry and kinematics with a physical energy-balance ansatz—each level relying on independent physical arguments:

Theorem 1 (Kinematic Matching Criterion). *In a space-filling vascular tree embedded in $\mathbb{R}^d$ ($d \geq 2$), the critical Womersley number Wo_c separating the viscous-dominated regime (overdamped, no coherent wave propagation) from the wave-dominated regime (underdamped, coherent reflections) is the unique solution of*

$$\mathcal{Q}_{\text{exact}}^{-1}(Wo_c) = d - 1, \quad (1)$$

where $\mathcal{Q}_{\text{exact}}^{-1}(Wo) \equiv \text{Re}(M'_{10}(Wo)) / |\text{Im}(M'_{10}(Wo))|$ is computed from the full Womersley solution of the Navier-Stokes equations for oscillatory flow in a rigid cylindrical tube, and M'_{10} is the modified Womersley function involving Bessel functions J_0, J_1 . For mammalian vascular trees ($d = 3$): $Wo_c = 1.740$ (exact numerical root). The criterion applies to two distinct physical quantities—the longitudinal fluid admittance Y_L and the characteristic wave admittance $Y_c \propto \sqrt{Y_L}$ —yielding two separate thresholds whose non-coincidence constitutes the mathematical formulation of incommensurability.

In plain terms, the Womersley number Wo measures how strongly fluid inertia competes with viscous friction over a single heartbeat, and the theorem identifies one critical value Wo_c that acts as a switch. *Below* it, viscous friction absorbs the sideways “kick” that the flow receives where the parent vessel forks, the junction stays quiet, and the tree behaves as a network of steady resistors—the viscous regime in which Murray’s law applies. *Above* it, that sideways energy can no longer be absorbed locally and is instead thrown back as a reflected pressure wave, so the tree behaves as a network of pulsatile waveguides—the wave regime. The two cost channels studied in this paper, viscous dissipation and wave reflection, do not therefore switch on at the same point: each has its own threshold, and it is the *gap* between them that we formalize as “incommensurability” and develop in the sections that follow. The remainder of this proof simply makes the switch quantitative, locating Wo_c from first principles.

Proof. The derivation proceeds through three independent levels:

Level 1 (Euclidean Geometry). A space-filling binary tree in $\mathbb{R}^d$ must maximize the angular separation between daughter branches to avoid overlap of perfusion domains. The optimal cone packing on S^{d-1} yields the regular simplex configuration, imposing the bifurcation half-angle $\cos(\varphi/2) = 1/\sqrt{d}$, whence $\tan^2(\varphi/2) = d - 1$. This is a theorem of discrete geometry, independent of any fluid property.

Level 2 (Classical Kinematics). An incident pulsatile velocity $\mathbf{v}$ along the parent axis decomposes at the junction into longitudinal (propagating) and transverse (deflected) components: $v_{\parallel} = v \cos(\varphi/2)$, $v_{\perp} = v \sin(\varphi/2)$. The ratio of transverse to longitudinal kinetic energy is therefore:

$$\frac{E_{\perp}}{E_{\parallel}} = \tan^2(\varphi/2) = d - 1. \quad (2)$$

This is vector algebra, not fluid dynamics.

Level 3 (Evanescent Mode Absorption). At the bifurcation, the transverse velocity component $v_{\perp}$ cannot propagate in the daughter vessels, which accept only axial flow; it constitutes an evanescent mode confined to the junction region. (An *evanescent* mode is flow energy that stays localized at the fork and decays with distance instead of travelling onward into the daughter vessels—it has nowhere to propagate, yet by energy conservation it cannot simply vanish.) By energy conservation, this evanescent energy must either be absorbed by viscous dissipation or scattered into a backward-propagating (reflected) wave in the parent vessel.

The viscous absorption capacity is quantified by $\mathcal{Q}^{-1} \equiv \text{Re}(Y_L) / |\text{Im}(Y_L)|$, which is the ratio of time-averaged dissipated power to reactive power—the standard loss tangent. Because both

the transverse and longitudinal modes oscillate at the cardiac frequency ω , the energy balance is naturally expressed in terms of mean powers: $P_{\text{diss}} = \mathcal{Q}^{-1} P_{\parallel}$. Complete absorption of the evanescent energy requires $P_{\text{diss}} \geq P_{\perp}$, i.e., $\mathcal{Q}^{-1} \geq d - 1$. Three regimes emerge:

- $\mathcal{Q}^{-1} > d - 1$ ($\text{Wo} < \text{Wo}_c$): *Overdamped*. Viscous absorption exceeds geometric scattering; all transverse energy is dissipated locally. No coherent wave propagation.
- $\mathcal{Q}^{-1} = d - 1$ ($\text{Wo} = \text{Wo}_c$): *Critical threshold*. Absorption exactly balances scattering—the onset of wave propagation.
- $\mathcal{Q}^{-1} < d - 1$ ($\text{Wo} > \text{Wo}_c$): *Underdamped*. Absorption is insufficient; excess transverse energy backscatters as a reflected wave ($\Gamma \neq 0$) at every junction.

The critical condition $\mathcal{Q}^{-1} = d - 1$ is the *vascular analogue of the Ioffe-Regel criterion*: the threshold where geometric scattering overwhelms viscous damping, in heuristic analogy with the onset of Anderson localization in disordered wave systems (a qualitative parallel invoked for physical intuition, not a formal correspondence). Since $\mathcal{Q}^{-1}(\text{Wo})$ is strictly monotonically decreasing, the solution is unique. $\square$

Corollary 2 (Closed-Form Approximation). *Using the low-Womersley asymptotic expansion of the Bessel functions ($1 - F_{10} \approx i\text{Wo}^2/8 + \text{Wo}^4/48$), the longitudinal admittance for pulsatile flow becomes:*

$$Y_L = Y_{\text{Poiseuille}} \times \left[1 - \frac{i\text{Wo}^2}{6} + \mathcal{O}(\text{Wo}^4) \right], \quad (3)$$

where $Y_{\text{Poiseuille}} = \pi r^4 / (8\mu\ell)$ is the steady-state Poiseuille conductance. The inverse quality factor is therefore:

$$\mathcal{Q}^{-1} = \frac{\text{Re}(Y_L)}{|\text{Im}(Y_L)|} = \frac{Y_{\text{Poiseuille}}}{Y_{\text{Poiseuille}} \cdot \text{Wo}^2/6} = \frac{6}{\text{Wo}^2} + \mathcal{O}(\text{Wo}^0). \quad (4)$$

Imposing the evanescent mode balance $\mathcal{Q}^{-1} = d - 1$ yields:

$$\text{Wo}_c^{\text{fluid}} \approx \sqrt{\frac{6}{d-1}}. \quad (5)$$

For $d = 3$: $\sqrt{3} \approx 1.732$ (relative error $< 0.5\%$ vs. the exact root 1.740). For $d = 2$: $\sqrt{6} \approx 2.449$.

Epistemological status. Each level employs an independent argument—discrete geometry (Level 1), vector kinematics (Level 2), and energy conservation with the Navier-Stokes admittance (Level 3). No level presupposes the result of another, eliminating circularity. We emphasise two technical points. First, the energy balance is expressed in terms of time-averaged powers, not energies per cycle; the definition $\mathcal{Q}^{-1} \equiv \text{Re}(Y)/|\text{Im}(Y)|$ directly gives the ratio of dissipated to reactive power, so no factor of 2π appears. Second, the evanescent-mode argument (Level 3) is a *physical ansatz*—a worst-case energy-balance hypothesis neglecting mode conversion to axial flow in daughter vessels—not a purely deductive mathematical necessity. This ansatz provides a sufficient condition for the transition threshold. The exact Bessel solution, which fully accounts for all modal coupling, confirms the threshold to within 0.5%, validating the physical hypothesis *post hoc*. The criterion is further validated by:

- Dimensional collapse: $d = 2$ retinal vasculature yields $\text{Wo}_c^{\text{fluid}} = \sqrt{6} \approx 2.449$, $\text{Wo}_c^{\text{wave}} = 0$ (Section 8),
- Allometric transition near $M^* \approx 0.84$ g separating the viscous and wave regimes.

This criterion establishes a correspondence between the *geometrical* degrees of freedom of the embedding space and the *dynamical* phase of the fluid. Levels 1–2 (topology and kinematics) are purely deductive; Level 3 (evanescent mode absorption) is a physical ansatz validated *post hoc* by the exact Bessel solution. The framework requires no phenomenological fitting. The minimax framework thus provides a metabolically invariant attractor requiring no species-specific tuning (Section 6).

2 The Information Cost of Local Optimization

The structural stability of vascular branching across seven orders of magnitude in body mass implies an optimization principle that is invariant to absolute physical scale. A fundamental question is whether this invariance can be implemented through purely local, junction-level feedback. Consider a local remodeling rule where a vessel adjusts its radius r_g to minimize a Lagrangian $\mathcal{L}_g = \mathcal{C}_{\text{met}} + \mu_g \mathcal{C}_{\text{wave}}$. Here, μ_g is a dimensionful coupling coefficient that weights the relative penalty of wave reflections against metabolic maintenance.

Proposition 3 (Symmetric Fine-Tuning Constraint for Local Rules). *Let $\mathcal{F}$ be a causal, locally-acting homeostatic rule mapping intensive and extensive physiological variables (pressure, shear stress, local impedance) to structural vessel adaptation. If the network is governed by two linearly independent physical dissipation regimes (e.g., viscous friction and reactive wave scattering), a perfectly symmetric local rule $\mathcal{F}$ cannot maintain a scale-invariant branching exponent α^* across the hierarchy without the local coupling parameters varying by orders of magnitude. Because such extreme fine-tuning is biologically implausible, real networks must break symmetry (e.g., via taper or asymmetric branching) to achieve scalable local optimization.*

Intuitively, the proposition says that a vessel acting on local information alone cannot keep the *entire* tree scale-free. Picture a bricklayer asked to build a self-similar pattern that must look the same whether the wall ends up one metre or one kilometre tall, but who can see only the brick in hand and its immediate neighbours. To place each brick correctly the bricklayer would need to know how far up the wall it sits—a global quantity. The vascular counterpart of “height up the wall” is the generation depth g , and the only way to infer it from local cues is through the absolute vessel size, which in turn requires knowing the size of the root (aortic) vessel. As Step 3 of the proof shows, that root scale is precisely the piece of information a downstream junction cannot recover from local measurements—which is why a strictly local, symmetric rule must fail.

Remark 2 (Biological Resolution via Structural Heterogeneity). Real vascular networks resolve this strict bound through structural heterogeneity: asymmetric branching ratios, diameter taper, and wall elasticity variations encode positional information implicitly, allowing the network to approximate scale-free behavior without explicit g -knowledge. The bound applies rigorously only to idealized symmetric networks.

Proof. We establish this constraint through four steps:

(1) Functional Independence. The gradients of the local power dissipation for Poiseuille flow and Womersley wave reflections scale with local absolute radius as $\mathcal{O}(r^{-5})$ and $\mathcal{O}(r^{-3})$, respectively. Being functionally independent, no stationary linear combination of these local costs possesses a continuous set of roots. Any static combination yields a specific absolute target radius, not a continuous scale-free hierarchy.

(2) The Epistemic Constraint. To maintain a constant α^* over a continuous domain of radii, $\mathcal{F}$ must dynamically adjust its internal coupling weights to perfectly cancel the shifting physical ratio $\Lambda(r) = \mathcal{C}_{\text{wave}}/\mathcal{C}_{\text{visc}} \propto r^2$. This compensation requires the explicit local computation of the

topological depth g . This r^2 growth is a direct consequence of Step (1): a viscous dissipation gradient scaling as $\mathcal{O}(r^{-5})$ integrates to a cost $\mathcal{C}_{\text{visc}} \propto r^{-4}$, while the $\mathcal{O}(r^{-3})$ wave gradient integrates to $\mathcal{C}_{\text{wave}} \propto r^{-2}$; their ratio therefore grows as $\Lambda(r) = \mathcal{C}_{\text{wave}}/\mathcal{C}_{\text{visc}} \propto r^{-2}/r^{-4} = r^2$, so a vessel ten times wider sees the wave penalty inflated a hundredfold relative to the viscous one.

(3) Information Asymmetry. Computing g from local geometric parameters requires inverting the network scaling law $g \propto \log(r_0/r_g)$. This computation strictly requires knowledge of r_0 (the root absolute scale). While downstream reflections encode the local input impedance, the physical information regarding the upstream root scale r_0 is topologically shielded and directionally inaccessible to a local junction. Like a transmission line terminated by a complex load, the local impedance measurement cannot uniquely determine the source characteristics.

(4) Conclusion. In a strictly symmetric tree, $\mathcal{F}$ is structurally bounded: it cannot access the non-local information (r_0) required to compute g , which is necessary to actively compensate $\Lambda(r)$. Consequently, maintaining a universal α^* via symmetric local gradient descent requires implausible local parameter fine-tuning. This implies that either the universal exponent is the scale-free attractor of a global network optimization principle, or biological networks must employ structural heterogeneities to circumvent the symmetric fine-tuning constraint. $\square$

2.1 Bio-Physical Limits of Endothelial Mechanotransduction

The fine-tuning constraint established by Proposition 3 for symmetric local rules is not merely a mathematical curiosity, but is strictly enforced by the biophysical and temporal constraints of the endothelial mechanosensory machinery:

1. **The Phase-Lag Blind Spot (Temporal Bound):** Endothelial cells respond to mechanical forces through complex signaling networks (e.g., Piezo1 mechanosensitive channels, integrins, and shear-induced phosphorylation cascades) which integrate signals over characteristic cellular time constants $\tau_{\text{bio}} \sim 100\text{--}1000$ ms. Given that mammalian cardiac periods span $T \sim 100\text{--}3000$ ms across the allometric range ($f_H \sim 0.3\text{--}10$ Hz, from large cetaceans to small rodents), and that cellular signaling time constants are comparable to or exceed these periods, these signaling networks operate as low-pass filters that are inherently blind to the sub-cycle phase-lag ϕ between pressure and velocity waves. Since evaluating the local inverse quality factor $\mathcal{Q}^{-1} = \text{Re}(Y)/|\text{Im}(Y)|$ (or the local reactance) requires resolving the precise cycle-by-cycle phase offset, the local cellular machinery is structurally incapable of measuring the wave transport parameters necessary for impedance-based tuning.
2. **Impedance Non-Locality (Spatial Bound):** The local reflection coefficient Γ_j at a vascular junction is determined by the input impedance of the entire downstream branching sub-tree. The mechanosensing apparatus of local endothelial cells (sensing only local Wall Shear Stress and circumferential wall tension) cannot decouple the chaotic superposition of backward reflections returning from millions of distal capillary junctions. Lacking a global feedback channel to probe this non-local topology, any local gradient descent algorithm is blind to the distal network's impedance state.
3. **The Murray Convergence Trap (Evolutionary Bound):** Homeostatic adaptation rules operating purely on local shear stress feedback (e.g., maintaining a constant shear stress target $\tau_w \approx c$) are mathematically guaranteed to converge to Murray's local attractor ($\alpha = 3$). However, in large conduit arteries, maintaining $\alpha = 3$ would cause catastrophic wave-reflection costs, severely compromising cardiac efficiency. The transition to $\alpha \approx 2.7$ in the large vessels of large mammals requires the cell to actively break the local Murray feedback loop. Because a cell cannot determine when to suppress this feedback without

knowing its global position g or the species-specific body mass M , the transition rule cannot be encoded locally, necessitating a global network-level minimax fitness landscape.

The network avoids this informational and sensory burden because the minimax attractor is an *evolutionary landscape constraint*. The cellular mechanosensors do not actively compute or solve the minimax optimization in real time; rather, the minimax operates as a physical attractor where mechanical fatigue and metabolic dissipation are globally minimized, shaping the genetic vascular layout through selective pressure.

Remark 3 (Genetic Encoding of Generation-Dependent Coupling μ_g). One could argue that the organism might bypass this local epistemic constraint by genetically hardcoding a generation-specific profile $\mu(g)$ directly into the biophysical remodeling machinery of the endothelial cells. However, this evolutionary strategy fails on two accounts: (i) **Non-Universality across Species**: Since the number of hierarchical vascular generations G scales with species body mass ($G \propto \log M$, ranging from $G \approx 10$ in small rodents to $G > 30$ in large cetaceans), a hardcoded genetic profile $\mu(g)$ would have to be completely reprogrammed for every species to maintain the universal exponent $\alpha^* \approx 2.72$. A change in the depth of the tree would shift the target exponent at each generation, destroying allometric scale-invariance. (ii) **Angiogenetic Growth Dynamics**: During angiogenesis, the vascular network expands dynamically by sprouting and adding new distal generations. A local vessel segment that is originally at generation g in a young animal would have to dynamically adjust its biophysical weight $\mu(g)$ to a new $\mu(g')$ as downstream segments grow, which requires real-time global coordination and feedback channels. Consequently, genetic hardcoding of a static $\mu(g)$ profile is developmentally and evolutionarily fragile, leaving the global network-level minimax attractor as the only robust, information-costless mechanism for scale-invariant transport.

Corollary 4 (Violation Conditions). *The Scaling Conflict Bound can be violated only if the organism implements non-local information channels that communicate the root scale r_0 to peripheral junctions. Mechanisms that would permit local optimization include: (i) persistent morphogen gradients encoding absolute position, (ii) retrograde chemical signaling from root to periphery, or (iii) centralized neural control with explicit g -encoding. Such mechanisms are not observed in mature mammalian vasculature, confirming the necessity of global optimization.*

Corollary 5 (Information-Accounting Argument). *Setting a generation-dependent profile that holds α^* constant requires each junction to know only its own generation index g —one of G levels, carrying at most $\log_2 G$ bits. The obstruction is therefore not the magnitude of this information (which is modest) but its non-local character: the generation index is a global label that cannot be reconstructed from purely local hemodynamic measurements.*

Proof. To regulate the wall profile $\mu(g)$ a junction need not resolve its full root-to-node address (which would cost $g \log_2 N$ bits); it requires only the scalar generation index g , i.e. $\log_2 G$ bits. However, a junction senses only local quantities (pressure, flow, wall shear) and hence only *ratios* such as $Z_{\text{local}}/Z_{\text{neighbour}}$, which depend on $r_g/r_{g\pm 1}$ but are invariant under a global rescaling of the tree. Determining the absolute level g (equivalently the radius ratio r_0/r_g) thus requires information that is absent from any local ratio; by the topological shielding of upstream signals (Proposition 3) it is also not transmitted from the root. The generation index is therefore informationally inaccessible to local control, *independently* of how few bits it nominally encodes—it is the non-locality, not a Shannon channel-capacity bound, that forbids local optimization. $\square$

The vascular tree avoids this informational burden by converging to a *network-level minimax* attractor. By shifting the optimization from the junction to the network hierarchy, the incommensurability problem is resolved without fine-tuning: universality emerges as a scale-free geometrical property, making the minimax an “information-costless” evolutionary strategy.

Remark 4 (The Heterogeneity Paradox). The extreme statistical heterogeneity reported in multi-study meta-analyses [5] is a direct signature of this local fragility. Vessels at different hierarchical levels inhabit different regimes in the (η, α) phase space; without a global minimax coordination, local samplings inevitably exhibit the deterministic spread observed empirically.

Remark 5 (Structural-Fluidic Shielding). As the hierarchy scales down ($Wo \rightarrow 0$), the rise of vessel wall thickness ($h/r \approx 0.42$) serves as a mechanical defense—a “Topological Shielding”—that preserves the transport map against pulsatile attenuation. This co-evolutionary strategy further suggests that the network optimizes for robust global invariants rather than unstable local couplings.

3 ATP Stoichiometry and Metabolic Scaling Symmetry

3.1 The Principle of Linear Energy Transduction

The choice of a cost functional for vascular optimization is not an arbitrary modeling decision but a consequence of the fundamental *ATP stoichiometry* of biological transport. In any system operating under metabolic pressure, every additional Watt of dissipated power $\Delta\Phi$ corresponds to a fixed, extensive quantity of glucose consumed per unit time. Because glucose is a fungible good with a constant caloric density, the fitness cost of vascular inefficiency is strictly proportional to the absolute energetic excess.

This physical constraint imposes a *Metabolic Scaling Symmetry*¹: the fitness penalty must be invariant under a uniform multiplicative rescaling of the basal metabolism, $\Phi_{\text{opt}} \rightarrow \lambda\Phi_{\text{opt}}$ where $\lambda \in \mathbb{R}^+$. To satisfy this invariance while maintaining allometric scale independence (i.e., a dimensionless penalty that allows a shrew and a whale to be governed by the same adaptive rules), the functional must take the form of a fractional excess. The linear penalty is thus not a modeling choice, but a consequence grounded in ATP stoichiometry and constrained by the first law of thermodynamics.

It is worth pausing on an apparent tension between the two statements just made—that the cost is proportional to the *absolute* energetic excess, yet must take the form of a *fractional* excess. The two are reconciled by separating *within*-organism from *across*-organism comparisons. At a fixed body size the penalty is linear in the absolute excess power, because ATP is charged watt by watt; but across body sizes spanning five orders of magnitude in absolute metabolism, selection can act only on the *dimensionless* excess, so the absolute power must be normalized by the organism’s own optimum Φ_{opt} . The fractional excess $(\Phi - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ is the unique form that is simultaneously linear in the wasted power *and* carries the same meaning for a shrew and a whale.

3.2 Thermodynamic Linearity and the Physical Occam’s Razor

The linear form of the cost functional $F = (\Phi - \Phi^*)/\Phi^*$ is not an arbitrary mathematical choice, but a strict consequence of oxidative phosphorylation stoichiometry and near-equilibrium thermodynamics. At the cellular level, excess fluidic dissipation is paid in units of ATP. Because this metabolic accounting is strictly additive—one extra Joule of dissipated physical work requires a strictly proportional increase in oxidative metabolism—the evolutionary fitness penalty must scale linearly with excess entropy production. Imposing non-linear penalties

¹Not a field-theoretic gauge symmetry (like U(1) or SU(2) in physics), but a global scaling invariance of the cost functional. In biological transport, this reflects the fact that fitness selection depends on relative metabolic excess, not absolute power consumption.

(e.g., quadratic or logarithmic) would imply unphysical non-stoichiometric metabolic costs, equivalent to postulating “compound interest” on ATP molecules.

Furthermore, numerical sensitivity analysis shows that imposing artificial non-linear penalties alters the theoretical attractor α^* by less than 0.3%. These alternative penalties are merely counterfactual numerical robustness tests, not physically admissible alternatives. Their negligible impact is the signature of Prigogine’s minimum entropy production principle in the Onsager regime: because the evolutionary attractor is strongly confining, the mature vascular network operates in the immediate vicinity of the fundamental physical lower bound ($\Phi \rightarrow \Phi^*$). In this near-optimal regime, the fitness penalty is governed by the excess dissipation $\Delta\Phi = \Phi - \Phi^*$, in which it is linear; since the optimum is a stationary point ($\nabla\Phi(\vec{x}^*) = 0$), $\Delta\Phi$ is in turn quadratic in the configurational deviation $\delta\vec{x}$.

Finally, the Gauge Invariance (normalization by the intrinsic absolute minimum Φ^*) represents the allometric “Occam’s razor”: it ensures that the selective pressure acting on a 30 g mouse is mathematically isomorphic to the pressure acting on a 3000 kg elephant, rendering the optimization topology scale-free and permitting universal scaling laws such as Kleiber’s metabolic rate to emerge naturally.

We formalize this thermodynamic necessity as a first-order response principle:

Theorem 6 (Onsager Linearity of Fitness Penalty). *Consider a vascular network characterized by a state vector $\vec{x} = (r_0, \dots, r_G, \alpha, \beta)$ describing vessel radii and branching exponents. Let $\vec{x}^*$ be the configuration minimizing total dissipation $\Phi(\vec{x})$ subject to morphometric constraints (fixed total volume and hierarchical connectivity).*

For small deviations $\delta\vec{x} = \vec{x} - \vec{x}^$ from the optimum, the fitness penalty $\mathcal{L}_{\text{fitness}}$ associated with increased metabolic cost is:*

$$\mathcal{L}_{\text{fitness}}(\vec{x}) = \lambda_{\text{bio}} \frac{\Phi(\vec{x}) - \Phi(\vec{x}^*)}{\Phi(\vec{x}^*)} + O(\|\delta\vec{x}\|^3) \quad (6)$$

where $\lambda_{\text{bio}} > 0$ is a phenotype-specific selection coefficient, and the penalty is linear in the excess dissipation $\Delta\Phi = \Phi - \Phi^*$; near the optimum $\Delta\Phi$ is itself $O(\|\delta\vec{x}\|^2)$, since $\nabla\Phi(\vec{x}^*) = 0$ (Step 1).

Moreover, the functional form $(\Phi - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ is the uniquely consistent scale-invariant measure of dissipative excess satisfying the joint requirements of scale invariance, thermodynamic extensivity, and normalization.

This theorem is the engine behind the linear cost functional used throughout the paper, so it is worth stating its content in words. It makes two linked claims. First, near the optimal architecture the evolutionary penalty grows in proportion to the *fractional* energy waste $\Delta\Phi/\Phi_{\text{opt}}$ —not its square, nor its logarithm—because metabolism charges a fixed number of ATP molecules per excess watt, with no “compound interest.” Second, that fractional waste is itself only a *second-order* (quadratic) function of how far the vessel geometry strays from optimal, because at a minimum the first derivative of the dissipation vanishes. Together these mean that a small geometric error incurs only a gentle, second-order metabolic penalty—which is exactly why real networks can tolerate the modest morphometric heterogeneity observed in practice while still sitting very close to the optimum.

Proof sketch; see Supplemental Material, Section S7 for full derivation. **Step 1: Near-equilibrium expansion.** For small deviations from the optimum:

$$\Phi(\vec{x}^* + \delta\vec{x}) = \Phi(\vec{x}^*) + \frac{1}{2} \delta\vec{x}^T \mathbf{H} \delta\vec{x} + O(\|\delta\vec{x}\|^3), \quad (7)$$

where $\mathbf{H}$ is the Hessian matrix (positive semi-definite at a minimum).

Step 2: Onsager entropy production. The excess entropy production rate is:

$$\Delta\dot{S} = \frac{\Delta\Phi}{T_{\text{body}}} \propto \|\delta\vec{x}\|^2, \quad (8)$$

quadratic in deviations (Onsager’s theorem for near-equilibrium systems [6]).

Step 3: Fractional fitness cost. Evolutionary selection acts on fractional metabolic increase, not absolute power (allometric invariance):

$$\mathcal{L}_{\text{fitness}} \sim \frac{\Delta\Phi}{\Phi_{\text{opt}}}. \quad (9)$$

The fractional form is necessary because absolute dissipation varies by 10^5 across species (mouse ~ 0.1 W, elephant $\sim 10^4$ W), but fitness selection is dimensionless.

Step 4: Uniqueness via Compositionality. The functional $(\Phi - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ is the simplest form satisfying:

- Scale invariance: $F(\Lambda\Phi, \Lambda\Phi_{\text{opt}}) = F(\Phi, \Phi_{\text{opt}})$,
- Normalization: $F(\Phi_{\text{opt}}, \Phi_{\text{opt}}) = 0$,
- Monotonicity: $\partial F/\partial\Phi > 0$.

While these axioms allow any monotonic $F(x/x_{\text{opt}})$, the additional requirement of **Compositionality**—that the optimality ranking of independent subsystems remains invariant under aggregation—singles out the linear (affine) form.

From scale invariance, F depends only on Φ/Φ_{opt} . Write $F(x) = f(x/1)$ where $x = \Phi/\Phi_{\text{opt}}$. For small $\epsilon = x - 1$:

$$f(1 + \epsilon) = f(1) + f'(1)\epsilon + O(\epsilon^2) = f'(1)\epsilon, \quad (10)$$

using $f(1) = 0$. Setting $f'(1) = 1$ (absorbing into λ_{bio}):

$$\mathcal{C}_{\text{met}} \propto \frac{\Phi - \Phi_{\text{opt}}}{\Phi_{\text{opt}}}. \quad (11)$$

□

Remark 6 (Robustness of Linearization in Physiological Regime). Empirical morphometry (Murray exponent $\alpha = 2.39 \pm 0.15$, meta-analysis [5]) suggests $\|\delta\vec{x}\|/\|\vec{x}^*\| \lesssim 0.1$, confirming operation in the linear response regime. Explicit calculation for the human vasculature at $\alpha = 2.39$ yields:

$$\frac{\Delta\Phi}{\Phi_{\text{opt}}} = \frac{\Phi(2.39) - \Phi(2.33)}{\Phi(2.33)} \approx 0.18\text{--}0.22, \quad (12)$$

where $\alpha_{\text{opt}} \approx 2.33$ for static transport. This fractional excess lies well within the regime where second-order corrections $O(\epsilon^2) \lesssim 0.05$ remain negligible, validating the linear approximation regardless of the specific functional form $(x - 1)$, $(x - 1)^2$, or $\ln x$.

While biological systems are fundamentally far-from-equilibrium, the Local Equilibrium Hypothesis (LEH) is well-justified for vascular remodeling, as the timescale of angiogenic adaptation is slow relative to the metabolic flux. This “angiogenic linearity” justifies the quadratic expansion of the excess entropy production, identifying the selection coefficient as a thermodynamic restoring force. $\lambda_{\text{bio}} \sim 0.2\text{--}0.3$ is consistent with cardiovascular traits under stabilizing selection.

3.3 Extensivity and Conditional Uniqueness of the Linear Form

Before establishing the full scale-invariant functional, we prove that the additivity (extensivity) of metabolic costs, when assumed, selects the linear penalty functional as the uniquely consistent form. This **Canonical Selection** provides a strong theoretical anchor for the exactness of the results, though alternative functionals are numerically near-degenerate: the logarithmic penalty shifts α^* by $\Delta\alpha^* \approx -0.004$, while the quadratic case yields $\Delta\alpha^* \approx +0.007$.

Biological justification for extensivity. In evolutionary biology, fitness costs associated with quantitative traits are commonly assumed to be additive when measured on fractional or logarithmic scales [7]. For metabolic traits, this reflects the stoichiometry of oxidative phosphorylation: each additional watt of excess dissipation costs a fixed number of ATP molecules, independent of the baseline metabolic flux. When two independent vascular networks A and B supply metabolically independent tissues, the total fitness penalty is the sum of the individual penalties, each normalized by its respective baseline consumption. Since the total baseline is $\Phi_{A,\text{opt}} + \Phi_{B,\text{opt}}$, the fractional excess of the combined system is naturally the weighted average

$$\mathcal{C}_{A+B} = w_A \mathcal{C}_A + w_B \mathcal{C}_B, \quad w_i = \frac{\Phi_{i,\text{opt}}}{\Phi_{A,\text{opt}} + \Phi_{B,\text{opt}}}. \quad (13)$$

This *extensibility* principle is an established framework in quantitative genetics and provides a thermodynamic foundation for the separability condition invoked below.

Compositionality and Segmentation Invariance. Beyond metabolic stoichiometry, the linear form is uniquely required by the requirement of **Segmentation Invariance**. If the vascular network is partitioned into arbitrary sub-networks (e.g., separating the arterial tree from the microvascular capillary bed), the total metabolic penalty must be independent of this arbitrary segmentation. As proven in Supplementary Section 7.1, any nonlinear penalty (e.g., quadratic or logarithmic) would introduce an artificial scaling dependence on the number of partitions, making the optimal exponent a function of the observer's chosen segmentation. The linear scaling form is the unique "Compositional Point" where the global optimization is scale-invariant under hierarchical partitioning.

Theorem 7 (Conditional Uniqueness of Linear Penalty Given Extensivity). *Let $\mathcal{C}(\alpha)$ be a dimensionless penalty functional for the total metabolic cost, defined on the extensive dissipation power $\Phi_{\text{net}}(\alpha)$. If the following extensivity condition holds:*

1. **Extensivity:** *For two independent vascular subsystems A and B with dissipations $\Phi_A(\alpha)$ and $\Phi_B(\alpha)$, the total penalty satisfies:*

$$\mathcal{C}_{A+B} = \frac{\Phi_A + \Phi_B - (\Phi_{A,\text{opt}} + \Phi_{B,\text{opt}})}{\Phi_{A,\text{opt}} + \Phi_{B,\text{opt}}} = \frac{\Phi_{A+B} - \Phi_{A+B,\text{opt}}}{\Phi_{A+B,\text{opt}}}. \quad (14)$$

2. **Separability:** *The penalty for the composite system can be expressed in terms of the individual penalties and their weights:*

$$\mathcal{C}_{A+B}(\alpha) = w_A \mathcal{C}_A(\alpha) + w_B \mathcal{C}_B(\alpha), \quad (15)$$

where $w_A = \Phi_{A,\text{opt}} / \Phi_{A+B,\text{opt}}$ and $w_B = \Phi_{B,\text{opt}} / \Phi_{A+B,\text{opt}}$.

Then $\mathcal{C}(\alpha)$ must be the affine (linear) functional:

$$\mathcal{C}(\alpha) = \frac{\Phi(\alpha) - \Phi_{\text{opt}}}{\Phi_{\text{opt}}}. \quad (16)$$

Proof. By the representation lemma (proven in the concluding step), scale invariance requires $\mathcal{C}(\alpha) = F(\Phi/\Phi_{\text{opt}})$ for some function F with $F(1) = 0$. Write $x = \Phi/\Phi_{\text{opt}}$.

From condition (2), for subsystems A and B :

$$F\left(\frac{\Phi_A + \Phi_B}{\Phi_{A,\text{opt}} + \Phi_{B,\text{opt}}}\right) = w_A F\left(\frac{\Phi_A}{\Phi_{A,\text{opt}}}\right) + w_B F\left(\frac{\Phi_B}{\Phi_{B,\text{opt}}}\right). \quad (17)$$

Substitute $\Phi_A = x_A \Phi_{A,\text{opt}}$ and $\Phi_B = x_B \Phi_{B,\text{opt}}$:

$$F(w_A x_A + w_B x_B) = w_A F(x_A) + w_B F(x_B). \quad (18)$$

In words, this equation demands that scoring the combined network directly (left-hand side) yield exactly the same number as averaging the two separate subsystem scores (right-hand side). A functional that failed it would rank a given vasculature as more or less wasteful depending only on where one happened to draw the boundary between subsystems, which is physically meaningless.

This is the *weighted Jensen (affinity) equation*: the condition obtained when Jensen's inequality is required to hold with equality for all weights w_A, w_B , not the classical midpoint Jensen equation $F\left(\frac{x+y}{2}\right) = \frac{F(x)+F(y)}{2}$. Assuming F is differentiable, differentiating both sides with respect to x_A yields:

$$w_A F'(w_A x_A + w_B x_B) = w_A F'(x_A). \quad (19)$$

Dividing by w_A gives $F'(w_A x_A + w_B x_B) = F'(x_A)$. Since this relation must hold independently of the value of x_B , the derivative F' must be a global constant c . Consequently, the second derivative is identically zero ($F''(x) = 0$), forcing F to be strictly affine:

$$F(x) = c(x - 1), \quad (20)$$

where the constant c is absorbed into the selection coefficient. Setting $c = 1$:

$$\mathcal{C}(\alpha) = \frac{\Phi(\alpha)}{\Phi_{\text{opt}}} - 1 = \frac{\Phi(\alpha) - \Phi_{\text{opt}}}{\Phi_{\text{opt}}}. \quad (21)$$

□

Remark 7 (Stochastic Expected Cost Interpretation). An alternative justification for the linear form emerges from a stochastic game-theoretic framework (Paper II [3]): if the vascular network faces uncertain metabolic loads with probability distribution $p(\eta)$, then minimizing the *expected* total cost $\mathbb{E}_\eta[\mathcal{C}_{\text{transport}} + \eta \mathcal{C}_{\text{wave}}]$ naturally leads to the linear combination of penalties. This provides a second, independent pathway to linearity beyond extensivity.

Remark 8 (Physical Interpretation). The extensivity condition reflects the fundamental additivity of metabolic costs: if two independent vascular networks each waste a fraction ϵ of their baseline metabolism, the combined system also wastes a fraction ϵ . This is a physical constraint from oxidative phosphorylation stoichiometry: each additional watt of excess dissipation costs a fixed number of ATP molecules, independent of the baseline flux.

This extensivity constraint strongly disfavors nonlinear alternatives:

- **Logarithmic penalty** $\ln(\Phi/\Phi_{\text{opt}})$: Violates extensivity due to Jensen's Inequality. Because the logarithmic function is strictly concave, the weighted average of the arguments strictly exceeds the weighted average of the logarithms: $\ln(w_A x_A + w_B x_B) > w_A \ln x_A + w_B \ln x_B$. This implies that the total penalty of the composite system grows superlinearly compared to the sum of its independent parts, introducing an unphysical compounding effect where larger aggregated systems are penalized disproportionately more than smaller subsystems. This violates the additive stoichiometry of ATP consumption, which scales linearly with absolute metabolic waste regardless of system size.

- **Quadratic penalty** $(\Phi/\Phi_{\text{opt}} - 1)^2$: Violates extensivity due to its strict convexity. By Jensen's Inequality, the quadratic penalty of a composite system with unequal subsystem performance ($\epsilon_A \neq \epsilon_B$) is always strictly lower than the weighted average of the individual penalties: $(w_A\epsilon_A + w_B\epsilon_B)^2 < w_A\epsilon_A^2 + w_B\epsilon_B^2$. This unphysically discounts the evolutionary penalty of unequal performance, violating the requirement that fitness costs scale additively with absolute metabolic waste.

The linear form emerges as the natural functional compatible with the thermodynamic requirement that fitness penalties scale with absolute metabolic waste. While the weighted-average extensivity condition is physically motivated by Gibbs additivity of free energy rather than rigorously derived from a variational principle, it provides a strong constraint that singles out the linear penalty as the simplest consistent form.

Physical basis and limitations. The extensivity condition reflects the stoichiometric additivity of ATP costs in oxidative phosphorylation and is consistent with quantitative genetics theory [7]. However, it is not mathematically derivable from first principles alone—it represents a biophysical hypothesis grounded in metabolic biochemistry. The uniqueness result (Theorem 7) is therefore *conditional* on this extensivity assumption. Alternative fitness landscapes (e.g., those involving resource competition or threshold effects) might violate extensivity, in which case the logarithmic or quadratic penalties could be biologically relevant. Empirical robustness tests (Supplemental Material, Section S7) show that alternative forms shift α^* by less than 0.3%, suggesting the linear form is effectively unique within the physiological regime.

3.4 Representation Theorem and Gauge Invariance

Lemma 8 (Representation Theorem for Gauge-Invariant Functionals). *Let $\mathcal{C}(\alpha)$ be a dimensionless functional of the branching exponent α defined on the network's total extensive transport power $\Phi_{\text{net}}(\alpha; \vec{\theta})$, where $\vec{\theta} = \{\mu_f, b, m_w, Q_0\}$ is the vector of scale and metabolic parameters. If $\mathcal{C}$ satisfies:*

1. **Gauge Invariance:** $\mathcal{C}(\alpha; \Lambda \vec{\theta}) = \mathcal{C}(\alpha; \vec{\theta})$ for all scale transformations $\Lambda > 0$,
2. **Ground State Normalization:** $\mathcal{C}(\alpha_{\text{opt}}) = 0$,

then $\mathcal{C}$ must have the form:

$$\mathcal{C}(\alpha) = F\left(\frac{\Phi_{\text{net}}(\alpha)}{\Phi_{\text{net}}(\alpha_{\text{opt}})}\right), \quad (22)$$

where $F: \mathbb{R}^+ \rightarrow \mathbb{R}$ is an arbitrary function with $F(1) = 0$.

Proof. By Euler's theorem for homogeneous functions, if $\Phi_{\text{net}}(\alpha; \vec{\theta})$ transforms under global scaling as $\Phi_{\text{net}}(\alpha; \Lambda \vec{\theta}) = f(\Lambda) \Phi_{\text{net}}(\alpha; \vec{\theta})$ for some function f , then scale invariance (condition 1) requires:

$$\mathcal{C}(\alpha; \Lambda \vec{\theta}) = \mathcal{C}(\alpha; \vec{\theta}). \quad (23)$$

By the Buckingham Π theorem, any scale-invariant quantity must be expressible as a function of dimensionless ratios formed from $\Phi_{\text{net}}(\alpha)$ and other quantities transforming identically under Λ . To avoid introducing arbitrary external parameters, the reference scale must be intrinsic to the network's phase space. The unique such reference state is the global transport minimum $\Phi_{\text{net}}(\alpha_{\text{opt}})$, which also transforms as $\Phi_{\text{opt}}(\Lambda \vec{\theta}) = f(\Lambda) \Phi_{\text{opt}}(\vec{\theta})$.

Therefore, $\mathcal{C}(\alpha)$ depends only on the dimensionless ratio $x = \Phi_{\text{net}}(\alpha) / \Phi_{\text{net}}(\alpha_{\text{opt}})$:

$$\mathcal{C}(\alpha) = F(x), \quad x = \frac{\Phi_{\text{net}}(\alpha)}{\Phi_{\text{net}}(\alpha_{\text{opt}})}. \quad (24)$$

Condition 2 (ground state normalization) requires $F(1) = 0$, completing the proof of Eq. (24) and establishing the representation lemma. $\square$

Theorem 9 (Scale-Free Normalization and Linearity of the Transport Penalty). *At fixed body mass M (and thus fixed heart rate $f_H \propto M^{-1/4}$), let $\Phi_{\text{net}}(\alpha; \vec{\theta})$ be the total extensive transport and metabolic power dissipated by the network. A valid dimensionless network-level penalty functional $\mathcal{C}_{\text{transport}}^{\text{net}}(\alpha)$ must satisfy:*

1. **Gauge Invariance (Scale Independence):** $\mathcal{C}(\alpha; \Lambda \vec{\theta}) = \mathcal{C}(\alpha; \vec{\theta})$.
2. **Positivity and Ground State:** $\mathcal{C}(\alpha) \geq 0$, with $\mathcal{C}(\alpha_{\text{opt}}) = 0$.
3. **Thermodynamic Linearity (Axiom 3):** *The penalty is strictly linear in the fractional excess dissipation. This linearity is supported by three independent arguments: (i) extensivity of metabolic costs selects the linear form as the unique consistent functional (Theorem 7, conditional on the extensivity hypothesis); (ii) near-equilibrium thermodynamics (Onsager regime, Theorem 6); (iii) robustness in the physiological regime where $|\Delta\Phi/\Phi_{\text{opt}}| \sim 0.2$ makes higher-order corrections negligible (see Remark following Theorem 6).*

The canonical functional satisfying these axioms is the fractional excess:

$$\mathcal{C}_{\text{transport}}^{\text{net}}(\alpha) = \frac{\Phi_{\text{net}}(\alpha) - \Phi_{\text{net}}(\alpha_{\text{opt}})}{\Phi_{\text{net}}(\alpha_{\text{opt}})}. \quad (25)$$

Proof. Independence of the three axioms. Axioms 1 and 2 together permit any positive-definite function of the dimensionless ratio $x = \Phi_{\text{net}}(\alpha)/\Phi_{\text{net}}(\alpha_{\text{opt}})$ satisfying $F(1) = 0$, such as $(x - 1)^2$ or $\ln x$. Axiom 3 is therefore logically independent and strictly necessary to isolate the affine form.

Justification of linearity. The linear penalty is not axiomatically imposed but follows from either:

- (a) **Extensivity** (Theorem 7): if assumed, uniquely selects linear form;
- (b) **Onsager near-equilibrium regime** (Theorem 6): empirically verified for vascular remodeling timescales;
- (c) **Physiological near-optimality:** $|\Delta\alpha/\alpha| \sim 0.1$ ensures higher-order corrections $O(\epsilon^2) \lesssim 0.05$ remain negligible.

Since (b) and (c) are empirical facts, linearity is *conditionally necessary* rather than axiomatically imposed.

Derivation. By Euler's theorem for homogeneous functions and the Buckingham Π theorem, Axiom 1 requires $\mathcal{C}(\alpha) = F(\Phi_{\text{net}}(\alpha)/\Phi_{\text{ref}})$, where Φ_{ref} transforms identically under Λ . To avoid introducing arbitrary external parameters, Φ_{ref} must be an intrinsic property of the network's phase space. The unique such reference state is the global transport minimum $\Phi_{\text{net}}(\alpha_{\text{opt}})$.

By the representation lemma, $\mathcal{C}(\alpha) = F(x)$ with $x = \Phi/\Phi_{\text{opt}}$. Axiom 2 gives $F(1) = 0$ and $F(x) \geq 0$ for $x \geq 1$.

To determine the unique form of F , write $x = 1 + \Delta\Phi/\Phi_{\text{opt}}$. Axiom 3 (thermodynamic linearity) requires that the penalty be strictly linear in the extensive excess entropy production $\Delta\Phi$, meaning $\partial\mathcal{C}/\partial\Phi$ is constant. This forces F to be affine in x :

$$F(x) = a(x - 1) + b, \quad a, b \in \mathbb{R}. \quad (26)$$

Applying the ground state normalization $F(1) = 0$ gives $b = 0$. Requiring $F(x) > 0$ for $x > 1$ (positivity of penalty for excess dissipation) gives $a > 0$. Up to a positive multiplicative constant, the unique solution is:

$$F(x) = x - 1, \quad (27)$$

yielding Eq. (25).

Consistency established: The fractional excess $(\Phi_{\text{net}} - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ is the *uniquely consistent* functional satisfying the thermodynamic requirements of extensivity and scale-invariant selection.

Exclusion of the logarithm. The logarithmic map $\ln(\Phi/\Phi_{\text{opt}})$, for instance, introduces a concave profile that dampens the evolutionary penalty for large structural deviations, decoupling it from the linear stoichiometry of oxidative phosphorylation. $\square$

Remark 9 (Epistemological Status: Parsimony Assumption vs. First Principle). Metabolic scale invariance—requiring the cost functional to be unchanged under $\vec{\theta} \rightarrow \Lambda\vec{\theta}$ —is not a first principle but an **assumption of parsimony** justified by allometric universality. Empirical metabolic rate scales as $M^{3/4}$ across 10^5 -fold mass range [8]. Any violation (i.e., penalty functional depending on absolute θ) would require a biologically privileged reference scale (e.g., $M_{\text{human}} = 70$ kg) that is absent from the data. We therefore impose scale-freeness as the simplest hypothesis consistent with observation. Alternative formulations introducing absolute thresholds $\theta_0(\text{taxon})$ could be tested if future data reveal systematic species-specific deviations from $M^{3/4}$ scaling; current evidence shows no such structure.

Remark 10 (Heart Rate as External Control Parameter). The heart rate f_H is *not* part of the scaling group $\vec{\theta}$ because it enters the Womersley number $Wo = r\sqrt{2\pi f_H/\nu}$ through the pulsatile frequency, not as a metabolic cost factor. The allometric scaling $f_H \propto M^{-1/4}$ acts as an *external control parameter* that drives the system through the viscous-to-wave transition as body mass varies. This explains why α^* is universal *within* a given mass class (e.g., all 70 kg mammals) but varies *between* mass classes (shrew vs. elephant). The scale invariance ensures universality at fixed M ; the allometric transition arises from the M -dependence of f_H .

Remark 11. Theorem 9 establishes that $\mathcal{C}_{\text{transport}}^{\text{net}}$ is not a phenomenological ansatz but the *only* dimensionless functional consistent with the symmetry of the biological optimization problem. The ground state Φ_{opt} plays the role of a natural gauge: it is intrinsic to the network, transforms homogeneously with $\vec{\theta}$, and produces the unique invariant ratio. Paper III [4] rigorously derives the metabolic scaling exponent $\beta_M(\alpha, d) = d\alpha/(2d + \alpha)$ (the organismal metabolic exponent, not to be confused with the branching ratio β) from proximal dominance (Theorem 1), showing that scale invariance uniquely maps local branching geometry to global organismal scaling laws.

Remark 12 (Gauge Invariance vs. Dimensional Homogeneity). It is crucial to distinguish **Scale-Free Normalization** (Axiom 1) from simple dimensional homogeneity. Dimensional analysis (Buckingham Π theorem) merely requires that the functional depend on dimensionless combinations of variables—but it does *not* uniquely specify which combinations or how they enter.

Scale invariance is a *physical symmetry principle*: the cost functional must remain invariant under global rescaling transformations $\vec{\theta} \rightarrow \Lambda\vec{\theta}$ that leave the physiological state unchanged (e.g., doubling all metabolic rates while doubling all flow rates). This symmetry *constrains the functional form*: it forces the cost to depend only on ratios $\Phi(\alpha)/\Phi_{\text{opt}}$, not on their absolute magnitudes or arbitrary external scales.

Combined with the linearity axiom (Axiom 3), scale invariance uniquely determines the cost functional to be the fractional excess $(\Phi - \Phi_{\text{opt}})/\Phi_{\text{opt}}$, excluding all other dimensionless forms (e.g., logarithmic, quadratic). This is strictly stronger than dimensional analysis, which would permit any function $F(\Phi/\Phi_{\text{opt}})$.

Remark 13 (Gauge Invariance vs. Field Theory Gauge). The term “scale invariance” here refers to the freedom to choose the reference scale Φ_{opt} as the metabolic normalization—analogueous to fixing the zero of the electrostatic potential in classical electrodynamics, not to local gauge symmetries (U(1), SU(2)) of field theory. The invariance $\mathcal{C}(\Lambda\vec{\theta}) = \mathcal{C}(\vec{\theta})$ is a *global scaling*

symmetry that, combined with normalization $F(1) = 0$, uniquely determines the functional form—precisely the role of gauge-fixing in physics. This terminology emphasizes that the cost functional depends only on dimensionless ratios, not on arbitrary external scales.

Remark 14 (Non-substitutability and aggregation robustness). The linear combination in $\mathcal{L}_{\text{net}}$ is the unique aggregation operator consistent with the physical non-substitutability of the two cost modes: wave energy loss and metabolic excess are orthogonal failure modes that cannot compensate each other. A multiplicative combination $\mathcal{L} \propto \mathcal{C}_{\text{wave}}^{\text{net}} \cdot \mathcal{C}_{\text{transport}}^{\text{net}}$ would permit arbitrarily large wave dissipation to be offset by arbitrarily efficient transport—a biologically inadmissible trade-off in a system where signal propagation and metabolic supply are independently required for viability. The minimax of a linear combination is the canonical robust-optimization formulation for non-fungible constraints [9], and the L_q robustness analysis of Paper II [3] confirms that the saddle point $\alpha^* = 2.72$ is stable across all convex aggregations $q \geq 1$, with a total spread $\Delta\alpha^* = 0.008$. A systematic test of the most natural alternative functional forms—quadratic, logarithmic, power-law, and absolute-threshold—against scale invariance, compositional consistency, Onsager linearity, and non-substitutability, confirming that the linear excess is the unique admissible choice, is provided in the Supplemental Material (Section S7).

Remark 15 (Third Commensurable Costs). One might ask whether the inclusion of a third physiological cost could render the optimization commensurable. The answer depends strictly on its physical dimensions. If the third cost is commensurable with metabolism (e.g., another energetic loss measured in watts), it is simply absorbed linearly into the transport gauge Φ_{net} , leaving the dimensionless structure of the Lagrangian unchanged. If it is incommensurable, it cannot linearly rescue the local optimization; rather, it introduces a third orthogonal axis to the minimax phase space, requiring a generalized zero-sum game over three competing invariants.

3.5 Architectural Decoupling: A Consistency Check

The scale invariance established above has a direct mathematical consequence: the minimax branching exponent α^* is determined entirely by dimensionless structural quantities and is *independent of every absolute metabolic parameter*. This architectural decoupling serves as a consistency check on the theory.

Corollary 10 (Structural Ground State and Scale-Free Normalization). *Let α^* be the minimax saddle point of the network Lagrangian $\mathcal{L}_{\text{net}}(\alpha, \eta)$ for an idealized symmetric, non-tapering vascular tree. Then α^* depends primarily on the dimensionless structural parameters $\mathcal{S} = (G, N, p, \alpha_w, A)$ and exhibits the following scaling behavior:*

$$\alpha^* = \alpha_{\text{ground}}^*(\mathcal{S}) + \delta\alpha_{\text{pert}}^*(\mu_f, \rho, \dots) \quad (28)$$

Pure metabolic parameters ($b, Q_0, \ell_0, \Delta G_{\text{ATP}}$) exhibit scale invariance: $|S_{\theta_i}| < 0.01$ (Table 1), confirming that the architecture is decoupled from absolute metabolic scales.

Fluid-mechanical parameters (μ_f, ρ) show moderate sensitivity ($|S_x| \approx 0.15\text{--}0.20$), reflecting their role in defining boundary conditions and the viscous-wave transition regime. The predicted ground state $\alpha_{\text{ground}}^* \approx 2.629$ represents the idealized symmetric attractor; structural heterogeneities (asymmetry, taper, elastic wall compliance) introduce perturbations $\delta\alpha^* \approx +0.091$, shifting large-mammal values toward the observed $\alpha^* \approx 2.72$.

Proof. By Theorem 9, the transport cost $\mathcal{C}_{\text{transport}}^{\text{net}}(\alpha)$ is degree-zero homogeneous in pure metabolic parameters $\vec{\theta}_{\text{met}}$, which cancel identically in the ratio $(\Phi_{\text{net}} - \Phi_{\text{opt}})/\Phi_{\text{opt}}$. The wave cost $\mathcal{C}_{\text{wave}}^{\text{net}}(\alpha)$ depends on the Womersley number $\text{Wo} \propto r\sqrt{2\pi f_H/\nu}$ (where $\nu = \mu_f/\rho$ is kinematic viscosity) and geometric branching ratios determined by $\mathcal{S}$. Changes in viscosity shift the viscous-wave

transition boundary, introducing moderate sensitivity to the minimax balance point. The ground state $\alpha_{\text{ground}}^*(\mathcal{S})$ is the solution for the idealized symmetric model; real biological systems exhibit perturbative shifts due to morphometric heterogeneity and fluid-mechanical variations. $\square$

Remark 16 (Numerical Verification). Table 1 confirms the two-regime structure: pure metabolic parameters show $|S_x| < 0.01$ (scale invariance), while fluid-mechanical and structural parameters show moderate sensitivity ($0.04 \leq |S_x| \leq 1.2$) across all mammalian species, confirming the robustness of the minimax attractor.

Remark 17 (Physical Interpretation). The Structural Ground State theorem establishes that the branching exponent is an *architectural attractor* rather than a rigidly fixed eigenvalue. The ground state $\alpha_{\text{ground}}^* \approx 2.629$ emerges from the network topology and dimensionless structural parameters, while moderate perturbations from fluid-mechanical variations and morphometric heterogeneities shift real systems toward $\alpha^* \approx 2.72$. This two-scale structure—a robust topological attractor with physiological perturbations—explains why the same approximate value appears across species whose metabolic rates and cardiac outputs differ by orders of magnitude, while still allowing for the subtle allometric trends observed in empirical data. The architecture is not *tuned* to the physiology; it is *attracted* to a structural ground state and modulated by biological perturbations.

4 Architectural Invariance of the Duty Cycle

Mathematical foundation: Robust Optimization. A Lagrangian structure of the form $\mathcal{L}_{\text{net}}(\alpha, \eta) = \eta \mathcal{C}_{\text{wave}} + (1 - \eta) \mathcal{C}_{\text{transport}}$ with $\eta \in [0, 1]$ corresponds, in the language of robust control theory [9], to a *Minimax Fractional Excess* problem: one minimizes, with respect to α , the worst-case fractional excess among incommensurable cost channels. The selection of η as a linear weight is the dual representation of a *worst-case* scenario uncertainty, where Nature selects the regime (metabolic or pulsatile) that produces the maximum penalty. The saddle point (α^*, η^*) satisfies the cost-balancing condition, which is the analog of the robust optimality condition for non-fungible constraints. Hereafter, we use “minimax” as a synonym for this construction, consistent with the terminology of robust programming.

It helps to read this construction as a game. The network chooses the branching geometry α ; an adversary—standing in for the uncontrolled physiological state (rest versus exercise, growth, a changing heart rate)—then chooses the weight η that makes the chosen geometry look as costly as possible, by tilting the balance toward whichever cost channel that geometry handles worst. A tree that is excellent for steady transport but reflects pulse waves can be punished by an adversary who emphasizes the wave channel, and a tree tuned to suppress reflections can be punished by one who emphasizes transport. The network’s safest move is therefore the geometry at which the adversary’s choice no longer matters—where the two costs are *exactly equal*, so tilting η in either direction changes nothing. That equal-cost geometry is the saddle point (α^*, η^*) , and its defining property is *immunity* to the operating weight: the architecture is optimized not for any single regime but against the worst case across all of them. This is the precise sense in which α^* is robust rather than fine-tuned.

Theorem 11 (Architectural Invariance of the Biological Duty Cycle). *Let (α^*, η^*) be the unique saddle point of the unified zero-sum network Lagrangian:*

$$\mathcal{L}_{\text{net}}(\alpha, \eta) = \eta \mathcal{C}_{\text{wave}}^{\text{net}}(\alpha) + (1 - \eta) \mathcal{C}_{\text{transport}}^{\text{net}}(\alpha). \quad (29)$$

The emergent marginal-balance weight η^ is an exact invariant of the network’s allometric class $\mathcal{A}(G, p, \alpha_w)$, strictly orthogonal to the absolute scale of the extensive metabolic parameters $\Lambda \vec{\theta}$.*

Proof. The Lagrangian $\mathcal{L}_{\text{net}}(\alpha, \eta)$ is affine (hence both convex and concave) in η and convex in α for every fixed $\eta \in [0, 1]$: $\mathcal{C}_{\text{transport}}$ is strictly convex, while $\mathcal{C}_{\text{wave}}$ has $\partial^2 \mathcal{C}_{\text{wave}} / \partial \alpha^2 > 0$ throughout the physiological range (verified below), hence is convex there, with a unique minimum at α_w . The domains $[\alpha_w, \alpha_t]$ and $[0, 1]$ are compact, and $\mathcal{L}_{\text{net}}$ is continuous. Because $\mathcal{L}_{\text{net}}$ is linear in η , the worst-case cost is attained at an endpoint,

$$\max_{\eta \in [0, 1]} \mathcal{L}_{\text{net}}(\alpha, \eta) = \max(\mathcal{C}_{\text{wave}}(\alpha), \mathcal{C}_{\text{transport}}(\alpha)),$$

the pointwise maximum of two convex functions and hence itself convex in α , so its minimum over the compact interval $[\alpha_w, \alpha_t]$ exists. Since $\mathcal{C}_{\text{wave}}$ increases away from α_w whereas $\mathcal{C}_{\text{transport}}$ decreases toward α_t on $[\alpha_w, \alpha_t]$, and the two costs cross exactly once in the interior, this minimiser is unique and satisfies the equal-cost condition $\mathcal{C}_{\text{wave}}(\alpha^*) = \mathcal{C}_{\text{transport}}(\alpha^*)$, giving the saddle point (α^*, η^*) . The affine-in- η , convex-in- α structure thus reduces existence and uniqueness of the saddle to elementary convex analysis, with no appeal to the general quasiconvex–quasiconcave minimax theorem. Detailed numerical verification (Supplemental Material, Table S1) confirms that the scalar Hessian $\partial^2 \mathcal{C}_{\text{wave}} / \partial \alpha^2 > 0$ throughout the physiological range $\alpha \in [2.0, 3.5]$ and $\eta \in [0, 1]$, securing the convexity on which this argument rests.

Uniqueness follows from the **strict convexity** of $\mathcal{C}_{\text{transport}}$ and the **strict monotonicity** of the marginal penalties: because $\partial \mathcal{C}_{\text{transport}}^{\text{net}} / \partial \alpha$ decreases monotonically from zero (becomes more negative) as α moves below α_t , and $\partial \mathcal{C}_{\text{wave}}^{\text{net}} / \partial \alpha$ increases monotonically from zero (becomes more positive) as α moves above α_w , their intersection at any fixed η is unique. The saddle-point condition $\partial \mathcal{L}_{\text{net}} / \partial \alpha = 0$ yields:

$$\eta^* = \left(1 + \frac{\partial \mathcal{C}_{\text{wave}}^{\text{net}}(\alpha^*) / \partial \alpha}{|\partial \mathcal{C}_{\text{transport}}^{\text{net}}(\alpha^*) / \partial \alpha|} \right)^{-1}. \quad (30)$$

Because $\mathcal{L}_{\text{net}}$ is affine in η , an interior saddle with $\eta^* \in (0, 1)$ can occur only when the maximiser is indifferent between the two endpoint costs; equivalently,

$$\frac{\partial \mathcal{L}_{\text{net}}}{\partial \eta} = \mathcal{C}_{\text{wave}}^{\text{net}}(\alpha^*) - \mathcal{C}_{\text{transport}}^{\text{net}}(\alpha^*) = 0,$$

which is the equal-cost condition. This is the formal counterpart of the game intuition above: at the saddle the adversary is indifferent to η precisely because the two costs coincide. This equivalence ensures that the emergent marginal-balance weight η^* effectively balances the two incommensurable penalties. The numerator derives from the acoustic impedance mismatch at bifurcations and is a strictly geometric function of $(G, N, \alpha_w, \alpha^*)$ containing no metabolic parameters.

By Theorem 9, $\mathcal{C}_{\text{transport}}^{\text{net}}$ is a degree-zero homogeneous function of $\vec{\theta}$. Its derivative with respect to α is therefore identically gauge-invariant:

$$\frac{\partial \mathcal{C}_{\text{transport}}^{\text{net}}}{\partial \alpha}(\alpha; \Lambda \vec{\theta}) = \frac{\partial \mathcal{C}_{\text{transport}}^{\text{net}}}{\partial \alpha}(\alpha; \vec{\theta}). \quad (31)$$

In Eq. (30) the watts cancel identically in the gradient ratio. The marginal-balance weight η^* is therefore a pure structural invariant, determined solely by (G, p, α_w) through the geometric evaluation of both derivatives at α^* . $\square$

Remark 18. This provides a first-principles explanation for the empirical observation that vascular branching exponents are conserved across developmental stages within a species, as suggested by the generation-invariant morphometry of [10] and the developmental convergence reported in [11].

Remark 19 (Theoretical Bounds on the Static Transport Attractor). Paper I [2] establishes that the static transport attractor α_t must lie within strict theoretical bounds derived from three-term metabolic cost structure: $(5 + p)/2 \approx 2.89 < \alpha_t < 3$ (Theorem 3.2), independent of flow asymmetry. Empirical values—ranging from $\alpha^* \approx 2.68$ in rat pulmonary arteries ($M = 430.0$ g) [12] to $\alpha^* \approx 2.72$ in large mammals [10]—lie below this static bound precisely because pulsatile dynamics contribute an additional wave-reflection penalty missing from purely viscous optimization, demonstrating that the minimax solution is the inevitable physical attractor within these constraints.

Remark 20 (Convexity Verification: Physical and Numerical). Existence and uniqueness of the saddle in Theorem 11 require only that $\mathcal{C}_{\text{wave}}(\alpha)$ be convex in α for each fixed η . We establish this through the following physical and numerical arguments:

Physical argument. The wave-reflection cost measures impedance mismatch at vascular junctions. The wave attractor exponent $\alpha_w = 2$ emerges from global reflection minimization (Paper II [3]): at each bifurcation, the power reflection coefficient $R^2(\alpha) = [(\gamma(\alpha) - 1)/(\gamma(\alpha) + 1)]^2$ with $\gamma(\alpha) = N^{1-2/\alpha}$ vanishes uniquely at $\alpha = 2$, selecting this value as the dynamically stable attractor over the network hierarchy. Away from this geometric impedance-matching condition ($\alpha = \alpha_w = 2$), the reflection coefficient $|\Gamma(\alpha)|^2$ grows *monotonically*: larger α (narrower daughters) increases mismatch in one direction, smaller α (wider daughters) increases it in the opposite direction. This monotonic growth ensures no local minima can exist. The cumulative reflection penalty $\mathcal{C}_{\text{wave}}(\alpha) \propto \sum_g |\Gamma_g(\alpha)|^2$ inherits this unimodal structure. Crucially, viscous damping (attenuation factor $e^{-2\kappa L_g}$ in the wave propagation) and spectral averaging over cardiac harmonics eliminate coherent interference effects that might otherwise introduce secondary extrema via resonances.

Numerical verification. Detailed numerical evaluation of $\mathcal{C}_{\text{wave}}(\alpha)$ for the physiological parameter range $\alpha \in [2.0, 3.5]$ and $\eta \in [0, 1]$ confirms strict unimodality (Supplemental Material, Table S1). No spurious local minima or inflection points are observed. The transport cost $\mathcal{C}_{\text{transport}}(\alpha)$ is *strictly* convex (proven analytically in Paper II [3] via power-law form in viscous dissipation). Since $\mathcal{C}_{\text{wave}}$ is convex over the physiological range and $\mathcal{C}_{\text{transport}}$ is strictly convex, their non-negative combination $\mathcal{L}_{\text{net}}(\alpha, \eta)$ is convex in α , so the convex-envelope argument of Theorem 11 yields a unique saddle point. The minimax equal-cost condition $f(\alpha^*) = g(\alpha^*)$ is verified without free parameters (Paper III [4]): for the porcine coronary tree, solving this balance yields $\alpha^* = 2.77$, consistent with the morphometric average $\alpha_{\text{exp}} = 2.70 \pm 0.20$, confirming that the network-level Lagrangian balances measurable first-principles costs.

Remark 21 (Allometric vs. angiogenic invariance). The invariance of η^* established in Theorem 11 is an *allometric* invariance: it holds under rescaling of the metabolic parameters $\vec{\theta}$ at fixed network topology (G, N, α_w, p) . It does not apply to *angiogenic* changes in G (addition or pruning of vascular generations), which alter $\kappa_{\text{eff}}(G)$ and thereby shift α^* . The two types of invariance are physically and mathematically distinct; only the former is established here.

Remark 22 (Physical status of Φ_{opt}). The reference state $\Phi_{\text{net}}(\alpha_{\text{opt}})$ is the cost attained when every vessel independently sits at its locally optimal radius $r^*(Q_g)$ —a configuration achievable only if the network is allowed to violate the global self-similar constraint $r_g = r_0 N^{-g/\alpha}$. It is therefore a mathematical lower bound (a gauge baseline), not a physiologically realizable state of the intact tree. The fractional excess $\mathcal{C}_{\text{transport}}^{\text{net}} = (\Phi_{\text{net}} - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ measures the architectural cost of imposing global self-similarity, not a deviation from a physically accessible optimum.

Remark 23 (The Static Attractor). As demonstrated in Paper I [2], the explicit evaluation of the static transport optimum α_t —with the inclusion of the metabolic cost of the vascular wall ($\propto r^{1+p}$)—strictly breaks the universality of Murray’s law ($\alpha = 3$) and fixes the purely static attractor at $\alpha_t \approx 2.900 - 2.940$ for mammalian coronary networks. This value constitutes the

theoretical starting point of the present framework: the residual gap from $\alpha_t \approx 2.900$ to the empirical $\alpha \approx 2.72$ is a mathematical necessity that requires the inclusion of the dynamic, pulsatile minimax mechanism developed here.

5 Recovery of Single-Mechanism Limits

Corollary 12 (Recovery of Single-Mechanism Limits). *Let $\eta \in [0, 1]$ be the cost-mixing weight of the unified Lagrangian $\mathcal{L}_{\text{net}}(\alpha, \eta)$. Then:*

1. **Static limit** ($\eta \rightarrow 0$): *The wave penalty vanishes and $\mathcal{L}_{\text{net}} \rightarrow \mathcal{C}_{\text{transport}}^{\text{net}}(\alpha)$. The unique minimizer is $\alpha^* = \alpha_t \in [2.900, 2.940]$, recovering the result of Paper I [2] exactly.*
2. **Wave-dominated limit** ($\eta \rightarrow 1$): *The transport penalty vanishes and $\mathcal{L}_{\text{net}} \rightarrow \mathcal{C}_{\text{wave}}^{\text{net}}(\alpha)$. The unique minimizer is $\alpha^* = \alpha_w = 2$, the impedance-matching attractor of Paper II [3]. (Note: $\alpha_w = 2.000$ is the theoretical rigid-wall area-preserving limit for binary branching ($N = 2$); the elastic-wall-corrected value incorporating histological scaling ($h \propto r^p$, $p = 0.77$) is $\alpha_w = 2.115$, as used in Fig. 1.)*
3. **Minimax interior** ($\eta = \eta^*$): *Neither mechanism dominates. By Theorem 11, the saddle point (α^*, η^*) is the unique interior solution. The symmetric model yields $\alpha_{\text{model}}^* = 2.629$, in quantitative agreement with rat pulmonary arteries ($M = 430.0 \text{ g}$, $\alpha \approx 2.68$, Jiang et al. 1994 [12]). The higher value observed in large mammals ($\alpha^* = 2.72$, Kassab 1993 [10]) reflects scale-dependent morphometric heterogeneities (see §7.5).*

The unified framework therefore contains Papers I and II as degenerate boundary cases of a single variational principle.

Biologically, these two boundary limits are not abstract endpoints but physically realized regimes—which is precisely why the single-mechanism laws were each discovered first, in isolation. The static limit ($\eta \rightarrow 0$, $\alpha^* \rightarrow \alpha_t$) is the regime of the smallest vessels and of all non-pulsatile transport systems: in arterioles and capillaries the Womersley number is small, flow is effectively quasi-steady, wave reflection carries negligible cost, and minimizing viscous dissipation alone recovers Murray’s law ($\alpha \rightarrow 3$)—the same exponent obeyed by the strictly non-pulsatile xylem and airway trees. The wave-dominated limit ($\eta \rightarrow 1$, $\alpha^* \rightarrow \alpha_w$) is the regime of the largest conduit arteries, where the Womersley number is large, pulsatile inertia dominates, and the network must instead match acoustic impedance across each junction to avoid reflecting cardiac power back toward the heart, recovering the area-preserving law ($\alpha \rightarrow 2$). The interior minimax (α^*, η^*) is therefore the architecture of the mid-calibre pulsatile beds—the coronary, cerebral, and pulmonary arteries that morphometric studies actually sample—where neither cost is negligible and the network must balance both incommensurable penalties at once. The unified Lagrangian thus explains *why* Murray’s law and the area-preserving law each hold only within their own vascular territory, and why the universal mammalian exponent sits strictly between them.

Remark 24 (Baseline vs. Elastic Wave Attractor). The symmetric model prediction $\alpha_{\text{model}}^* = 2.629$ uses the baseline rigid-cylinder wave attractor $\alpha_w = 2.0$, corresponding to acoustic impedance matching $Z \propto d^{-2}$ in the absence of wall compliance. Paper II [3] incorporates elastic wall effects via the histological scaling $h \propto r^p$, yielding the modified wave attractor $\alpha_w = (5 - p)/2 \approx 2.115$ for $p = 0.77$. The elastic correction alone shifts the symmetric minimax to $\alpha^* \approx 2.655$; the full coherent impedance solver of Paper II—integrating multiple wave reflections, bifurcation asymmetry, vessel taper, and the Fåhræus-Lindqvist viscosity shift—recovers $\alpha^* \approx 2.72$ for porcine coronaries ($G = 11$, $N = 2$), in quantitative agreement with morphometric data [10].

The present work employs the baseline $\alpha_w = 2.0$ to isolate the fundamental incommensurability mechanism from material-specific corrections. The wall thickness exponent $p = 0.77$ is an independently measured empirical constant from Kassab’s morphometric data [10], not a free fitting parameter used to match the predicted α^* to observations.

Figure 1 illustrates the one-dimensional phase diagram of the unified Lagrangian, with the two boundary attractors and the unique robust interior saddle point.

Theorem 9 predicts that α^* is invariant to all absolute metabolic scales and sensitive only to structural parameters. Table 1 confirms this numerically on the porcine coronary tree ($G = 11$, $p = 0.77$, $\alpha_w = 2.000$, $N = 2$) using the generation-by-generation network-level Lagrangian, which yields $\alpha_{\text{model}}^* = 2.629$. The residual $\Delta\alpha^* = 0.091$ relative to the empirically calibrated value ($\alpha^* = 2.72$) reflects the simplified symmetric, non-tapering architecture used in the sensitivity model; the full morphometric model incorporating bifurcation asymmetry and vessel taper closes the gap (see §7.5).

Table 1: **Sensitivity analysis of α^* to physiological parameters.** Log-sensitivity $S_x = \partial\alpha^*/\partial \ln x$ computed numerically on the porcine coronary tree ($G = 11$, $p = 0.77$, $\alpha_w = 2.000$, $N = 2$) with the network-level cost functions of Paper II [3]. **Pure metabolic parameters** (blood cost b , proximal flow Q_0 , segment length ℓ_0) show $|S_x| < 0.01$, confirming metabolic gauge invariance (Theorem 9). These are precisely the metabolic quantities named in the abstract: the proximal flow Q_0 represents cardiac output, and the blood metabolic cost b encodes the ATP/oxygen stoichiometry; the table thus directly substantiates the independence claimed there. **Structural and fluid-mechanical parameters** (viscosity μ_f , wall metabolism m_w , wall exponent p , tree depth G , wave exponent α_w) show moderate sensitivity ($0.04 \leq |S_x| \leq 1.2$), reflecting their role in defining the boundary conditions and topology of the optimization problem. Despite this moderate parametric sensitivity, the predicted α^* remains confined to the narrow physiological window $[2.629, 2.72]$ across all mammalian species.

Parameter (x)	Baseline	Perturbation	$ S_x $
<i>Pure metabolic parameters (gauge-invariant: $S_x < 0.01$)</i>			
Blood cost (b)	1930.0 W m ⁻³	±10%	0.0082
Proximal flow (Q_0)	1.3 mL s ⁻¹	±10%	0.0106
Segment length (ℓ_0)	15.0 mm	±10%	0.0106
<i>Physiological & fluid-mechanical parameters (sensitive)</i>			
Wall metabolism (m_w)	20.0 kW m ⁻³	±10%	0.1820
Viscosity (μ_f)	3.5 mPa·s	±10%	0.1917
<i>Structural inputs (architecture-dependent)</i>			
Wall exponent (p)	0.77	±10%	1.1706
Tree depth (G)	11	±1	0.1803
Wave exponent (α_w)	2.000	±10%	0.0220




Figure 1: **One-dimensional phase diagram of the unified network Lagrangian.** The wave-impedance attractor ($\alpha_w = 2.115$, $\eta \rightarrow 1$) and the static transport attractor ($\alpha_t \approx 2.920$, $\eta \rightarrow 0$) are the two degenerate boundary cases of the unified Lagrangian $\mathcal{L}_{\text{net}}(\alpha, \eta)$. The physiological branching exponent $\alpha^* = 2.72$ emerges as the unique robust minimax saddle point at marginal-balance weight $\eta^* \approx 0.814$, stabilised by evolutionary selection pressure from both extremes. The cost-mixing weight η and the branching exponent α exhibit a monotonic inverse relationship at the saddle point: increasing α corresponds to a shift toward metabolic dominance (lower η).

6 The Architectural Transition

The gauge-invariant network Lagrangian established in §3 assumes a scale-invariant wave penalty $\mathcal{C}_{\text{wave}}$. This is exact only in the limit $\text{Wo} \rightarrow \infty$, where fluid inertia dominates and the impedance of a vessel depends solely on its cross-sectional area. In the general case, the reflection at each junction is governed by the mismatch of the complex Womersley admittances $Y(\text{Wo}) \propto R^2 \sqrt{1 - F_{10}(\text{Wo})}$.

Theorem 13 (Analytic Admittance Matching). *The universal vascular branching exponent α is bounded by two exact fluid-dynamic attractors representing the limits of zero reflection ($\Gamma = 0$):*

1. **Inertial Attractor** ($\text{Wo} \rightarrow \infty$): $F_{10} \rightarrow 0$, thus $Y \propto R^2$. Impedance matching requires $R_p^2 = \sum R_{c,i}^2$, yielding the area-preserving law $\alpha = 2$.
2. **Viscous Attractor** ($\text{Wo} \rightarrow 0$): $1 - F_{10} \propto \text{Wo}^2$, thus $\sqrt{1 - F_{10}} \propto \text{Wo}$. Since $\text{Wo} \propto R$, the admittance scales as $Y \propto R^3$. Impedance matching requires $R_p^3 = \sum R_{c,i}^3$, yielding Murray's Law $\alpha = 3$.

The allometric transition is thus the trajectory of the network-level minimax as the characteristic Womersley number scales with body mass M . To ensure a first-principles derivation, the wave penalty $\mathcal{C}_{\text{wave}}$ is defined as the spectral sum of reflected power across the first five harmonics $\{n\omega_H\}$ of the cardiac cycle, weighted by the longitudinal k -dispersion attenuation $e^{-2\kappa L}$ (see Supplemental Material, Section S2):

$$\mathcal{C}_{\text{wave}} = \sum_{n=1}^5 H_n \sum_{g=0}^{G-1} \left(\prod_{j=0}^g \eta_{\text{att},j}^{(n)} \right) |\Gamma_g^{(n)}|^2, \quad (32)$$

where H_n are the fixed power weights of the mammalian pulse spectrum. This formulation requires no ad-hoc damping functions or scaling constants. Sensitivity analysis (see Supplemental Material, Section S4) demonstrates that α^* is robust to variations in the spectral envelope, as the fundamental harmonic H_1 dominates the energetic partition. The ‘viscous shielding’ of

the microcirculation emerges naturally from the exponential decay of oscillatory energy in the low- Wo regime. The critical transition at $M^* \approx 0.84$ g is the physical threshold where viscous dissipation cedes dominance to the incommensurability principle. This threshold is governed by the scaling of the characteristic Womersley number:

$$Wo_g \approx Wo_{\text{ref}} \left(\frac{r_g}{r_{\text{ref}}} \right) \sqrt{\frac{f}{f_{\text{ref}}} \frac{\nu_{\text{ref}}}{\nu}}, \quad (33)$$

where Wo_{ref} is the value at a known physiological scale (e.g., human adult aorta).

Theorem 14 (The Cross-Class Allometric Transition). *While the minimax attractor α^* is structurally invariant within a fixed allometric class (Theorem 11), the transition between distinct regimes (viscous vs. wave-dominated) is governed by the crossing of a critical Womersley interval $[Wo_c^{\text{fluid}}, Wo_c^{\text{wave}}] = [\sqrt{3}, 3/\sqrt{2}]$. This dual-threshold transition induces a sharp allometric shift in branching geometry at a body mass $M^* \approx 0.84$ g.*

Theorem 15 (Emergent Allometric Transition Limits). *For a given body mass M , let $\alpha^*(M)$ be the unique minimax saddle point of $\mathcal{L}_{\text{net}}(\alpha, \eta; M)$. Then the trajectory of the attractor obeys the following limits:*

1. As $M \rightarrow 0$, $\alpha^*(M) \rightarrow \alpha_t \approx 2.920$ (viscous dominance).
2. As $M \rightarrow \infty$, $\alpha^*(M) \rightarrow \alpha_w \approx 2.0$ (inertial dominance).
3. There exists a unique mass $M^* \approx 0.84$ g where the slope $|d\alpha^*/d(\ln M)|$ is maximal, defining the sigmoidal inflection point of the allometric transition.

Proof. The limits follow from the asymptotic behavior of the admittance $Y(Wo)$. For $M \rightarrow 0$, $Wo \rightarrow 0$ throughout the tree, the reflection coefficient tends to zero, removing the wave penalty and allowing the transport cost to pull α^* toward its static minimum at $\alpha_t \approx 2.920$. For $M \rightarrow \infty$, the inertial geometric matching condition $\alpha_w \approx 2.0$ dominates. The sigmoidal transition in Fig. 2 emerges directly from this numerical competition, smoothly connecting the two topological attractors. $\square$

Corollary 16 (Fourth-Power Scaling of Transition Mass). *The critical mass M^* at which the allometric transition occurs scales as the fourth power of the critical Womersley number:*

$$M^* \propto Wo_c^4. \quad (34)$$

For networks embedded in different dimensions, this yields a topologically grounded prediction:

$$\frac{M_{d=2}^*}{M_{d=3}^*} = \left(\frac{Wo_c(d=2)}{Wo_c(d=3)} \right)^4 = \left(\frac{\sqrt{6}}{\sqrt{3}} \right)^4 = 4. \quad (35)$$

Proof. The transition mass is defined by the condition that the characteristic Womersley number at a reference generation (e.g., the aorta) reaches the critical fluid threshold Wo_c^{fluid} . From the allometric scaling of heart rate $f \propto M^{-1/4}$ and aortic radius $r_0 \propto M^{3/8}$ (derived from Kleiber's Law via flow continuity), the dimensional form of the Womersley number gives:

$$Wo_0 = r_0 \sqrt{\frac{2\pi f}{\nu}} \propto M^{3/8} \cdot M^{-1/8} = M^{1/4}. \quad (36)$$

Setting $Wo_0(M^*) = Wo_c$ yields the scaling relation $M^* \propto Wo_c^4$. However, this dimensional estimate (which would predict $M^* \sim 5\text{--}10$ g for typical mammalian parameters) represents only

an order-of-magnitude threshold. The precise value $M^* \approx 0.84 \text{ g}$ is determined numerically as the inflection point of the sigmoidal transition $\alpha^*(M)$, i.e., the body mass at which $|d\alpha^*/d(\ln M)|$ is maximal (Theorem 15, item 3). This sharper definition captures the smooth, progressive nature of the allometric crossover, rather than a discontinuous jump at the Womersley threshold.

The dimensional ratio (35) follows immediately from the kinematic matching criterion applied to the bulk fluid $\text{Wo}_c^{\text{fluid}}(d) = \sqrt{6/(d-1)}$ derived in §6.1. This prediction is experimentally testable by comparing the allometric transitions in planar networks (e.g., retinal vasculature) versus three-dimensional networks (e.g., coronary circulation). $\square$

6.1 Kinematic Matching Criterion at the Bifurcation

The architectural transition at M^* marks the critical threshold where the vascular tree can no longer ignore the complex phase of the fluidic impedance. To determine this critical Womersley number Wo_c , we must address a fundamental epistemological boundary in biological transport modeling.

The 1D Navier-Stokes equations governing pulsatile flow within a single vascular segment yield the complex Womersley admittance $Y(\text{Wo})$. However, these local equations of motion are *topologically blind*: they contain no information regarding the dimensionality d of the space the network is required to fill. Consequently, it is mathematically impossible to derive a global geometric threshold strictly from the local fluid dynamics. For the 1D fluidic transport to stably support a space-filling hierarchy, the local hydrodynamics must geometrically couple with the d -dimensional structural embedding.

Rather than relying on phenomenological fitting, we resolve this mechanical-topological decoupling by applying the **Kinematic Matching Criterion** (Theorem 1). This criterion establishes a rigorous condition—derived from Euclidean geometry, classical kinematics, and evanescent mode absorption at the junction—demonstrating that the pulsatile network undergoes a critical transition when the viscous absorption capacity of the Womersley flow (quantified by the inverse quality factor $\mathcal{Q}^{-1}$) exactly balances the geometric scattering imposed by the space-filling topology. For the network to maintain this equilibrium, the inverse quality factor must satisfy:

$$\mathcal{Q}^{-1} = d - 1. \quad (37)$$

This criterion acts as the essential boundary condition bridging 1D transport and d -dimensional topology. For mammalian arterial beds embedded in 3D tissue ($d = 3$), this requires $\mathcal{Q}^{-1} = 2$. By combining the Navier-Stokes expression for viscous absorption ($\mathcal{Q}^{-1} = 6/\text{Wo}^2$) with the Euclidean space-filling geometry, the criterion provides the first-principles foundation for the allometric transition.

However, pulsatile transport involves two mathematically distinct physical quantities: the *local fluid mass* (governed by the longitudinal admittance Y_L) which undergoes viscous dissipation, and the *propagating pressure wave* (governed by the characteristic admittance Y_c) which dictates network reflections. Because $Y_c \propto \sqrt{Y_L}$, the wave phase is algebraically halved with respect to the bulk fluid, inducing a profound **Phase Decoupling**. Applying the kinematic matching criterion ($\mathcal{Q}^{-1} = d - 1$) to both quantities yields two distinct thresholds (derived in Supplemental Section S1). We emphasise that the isotropic-embedding criterion $\mathcal{Q}^{-1} = d - 1$ is derived in S1 for the longitudinal (bulk-fluid) mode; its application to the characteristic admittance Y_c is a *postulate*—we extend the same kinematic matching to the propagating wave by analogy, with the algebraic phase-halving $Y_c \propto \sqrt{Y_L}$ as its physical basis. We adopt this as a working hypothesis rather than a derived result; the qualitative conclusion (two incommensurable thresholds) follows from the phase-halving alone and is robust to the precise value of the wave threshold.

The intuition behind the two thresholds is that a bifurcation must cope with *two* things at once: the bulk fluid it carries, and the pressure wave riding on that fluid. As an animal grows and its Womersley number rises, the bulk fluid is the first to “switch on”—to stop dissipating its junction kick viscously and begin reflecting it—at Wo_c^{fluid} . The pressure wave, because it responds to the *square root* of the fluid admittance, lags behind and switches on only at the larger Wo_c^{wave} . The two can therefore never be perfectly matched at the same body size: whatever the animal’s scale, one channel is always on the “wrong” side of its own threshold. That permanent, unavoidable mismatch is precisely what we mean by incommensurability, and it is why no single branching exponent can zero out both costs at once.

1. **The Fluid Threshold (Y_L):** The local bulk fluid achieves space-filling resonance when its dissipation ratio matches the topology:

$$\frac{6}{Wo^2} = d - 1 \quad \Rightarrow \quad Wo_c^{\text{fluid}} = \sqrt{\frac{6}{d-1}}. \quad (38)$$

For $d = 3$, this yields $Wo_c^{\text{fluid}} = \sqrt{3} \approx 1.732$.

2. **The Wave Threshold (Y_c):** The propagating wave, subjected to the square root of the local admittance, achieves kinematic matching at a higher Womersley number:

$$\frac{\sqrt{36 + Wo^4} + Wo^2}{6} = d - 1 \quad \Rightarrow \quad Wo_c^{\text{wave}} = \sqrt{\frac{3d(d-2)}{d-1}}. \quad (39)$$

For $d = 3$, this yields $Wo_c^{\text{wave}} = \frac{3}{\sqrt{2}} \approx 2.121$.

The fact that these two thresholds do not coincide constitutes the exact mathematical formulation of the **Incommensurability Principle**: the fluid and the wave can never achieve equipartition simultaneously. The ratio of their squared thresholds, $\frac{d(d-2)}{2}$, is never equal to 1 for any integer d . Remarkably, Womersley’s historical empirical transition threshold ($Wo_c \approx 2.0$, empirically utilized in Paper III) falls precisely within the theoretical gap $[\sqrt{3}, 3/\sqrt{2}]$ defined by this decoupling, perfectly bridging the phenomenological observations with first-principles theory.

It should be stated plainly that the *quantitative* predictions of this work—the transition mass M^* , the $\alpha^*(M)$ curve, and the per-species exponents—enter the solver only through the fluid threshold $Wo_c^{\text{fluid}} = \sqrt{3}$, which sets the viscous-to-inertial crossover used to compute $\alpha^*(M)$. The wave threshold $Wo_c^{\text{wave}} = 3/\sqrt{2}$ fixes the *upper* edge of the incommensurability gap and the qualitative regime structure (and the $d = 2$ retinal collapse), but it does not appear in the numerical optimization. The incommensurability of the two thresholds is thus a structural, regime-defining statement; the headline numbers depend only on $\sqrt{3}$.

The predictive power of this Phase Decoupling is confirmed by its global consequences. First, it defines a **scaling transition interval** in body mass $M^* \sim 1$ g (plausible range 0.5–1.2 g depending on species-specific heart rate variability), separating viscous-dominated from wave-influenced mammalian architectures. Second, it provides a definitive structural explanation for the “Retinal Paradox” (Mechanical-Fluidic Decoupling, §8): for networks constrained to a 2D manifold (e.g., the human retina, $d = 2$), the wave threshold collapses exactly to zero: $Wo_c^{\text{wave}}(d=2) = 0$. This implies that in a 2D topology, the propagating wave is *permanently* above its threshold ($Q_{Y_c}^{-1} \geq 1$ for all $Wo > 0$). Here the criterion is applied to the characteristic (transmitted) admittance Y_c , for which a large $Q_{Y_c}^{-1}$ denotes a permanently impedance-matched, propagating wave channel—the opposite sense to the longitudinal-mode damping convention, in which a large Q^{-1} signals overdamping; the two senses must not be conflated. This geometric

constraint forces planar retinal diameters to permanently transition toward the wave-attractor limit ($\alpha \approx 2.0$), directly confirming that the branching geometry is topologically anchored by the embedding dimension rather than local metabolic fine-tuning.

Falsifiable Milestones. The theory predicts that organisms below $M^* \approx 0.84$ g (sub-gram invertebrates, early neonates) should exhibit branching exponents closer to the Murray attractor ($\alpha \approx 3$), while all vascular mammals above this threshold should share the wave-influenced attractor ($\alpha \approx 2.5$ – 2.8). This prediction is testable via morphometric analysis of vascular casts across developmental stages.

Table 2: **Womersley Transition scaling (β) vs Body Mass.** Comparison of the theoretical branching parameter β predicted by the emergent minimax framework. The theoretical model reproduces the architectural transition from viscous-dominated morphologies ($\beta \approx 0.85$) to the pulsatile attractor ($\beta \approx 0.77$) without free parameters. The simplified model predicts a nearly constant branching ratio $\beta \approx 0.77$ for all mammals; the observed variation (see, e.g., [10, 13]) requires the inclusion of taper and asymmetry as in Paper II.

Species	Mass (M)	β (Phenomenological)	β (Emergent Minimax)
Shrew	3.0 g	0.82	0.7807
Mouse	25.0 g	0.80	0.7767
Rat	430.0 g	0.79	0.7726
Guinea Pig	700.0 g	0.785	0.7720
Rabbit	3000.0 g	0.78	0.7706
Human	70000.0 g	0.77	0.7682
Horse	500000.0 g	0.765	0.7672




Figure 2: The Incommensurability Jump: Allometric Transition of Branching Geometry. The emergent branching exponent $\alpha^*(M)$ as a function of body mass, computed from the symmetric minimax saddle point with no free parameters. Horizontal dashed lines mark the two attractors: the transport limit $\alpha_t \approx 2.920$ ($\eta \rightarrow 0$) and the wave limit $\alpha_w = 2.000$ ($\eta \rightarrow 1$; this is the rigid-wall baseline used to compute the symmetric minimax curve, distinct from the elastic-corrected $\alpha_w = 2.115$ shown in the conceptual phase diagram, Fig. 1). The shaded band (± 0.10) represents the estimated range of morphometric heterogeneity corrections (bifurcation asymmetry, taper; see §7.5). The vertical dotted line marks the predicted transition mass $M^* \approx 0.84$ g. Quantitative per-species validation is provided in Table 2.

6.2 Meta-Analysis of the Allometric Transition

The transition observed in Figure 2 suggests that mammalian evolution is governed by a “metacritical” boundary at $M^* \approx 0.84$ g. This threshold is not merely a statistical centroid but corresponds to the biological transition from viscous-dominated transport to wave-dominated pulsatility.

The simplified model places the transition at $M^* \approx 0.84$ g, the boundary between sub-gram organisms and the smallest mammals. For all mammals above this threshold, the minimax attractor settles near $\alpha^* \approx 2.629$, with only weak residual mass dependence. The full morphometric analysis of Paper II, which incorporates generation-by-generation vessel taper and asymmetry, is required to shift the transition to the 30–60 g range where rodent weaning occurs; this remains a quantitative refinement for future work. We stress that the falsifiable predictions below are anchored on the Womersley number relative to $Wo_c = \sqrt{3}$, not on the absolute

transition mass: M^* is a model-dependent landmark—the single-scale estimate 0.84 g shifts upward to tens of grams once taper and asymmetry rescale the Wo -mass mapping—whereas the hummingbird and neonatal tests turn only on whether Wo_0 crosses $\sqrt{3}$ and are therefore robust to this refinement. **The Hummingbird Test: A Falsifiable Prediction.** A critical experimental test of the Wo -driven transition arises in high-performance avian physiology. Despite tiny body mass, extreme heart rate can push small organisms into the wave-dominated regime.

Table 3: Predictive verification of the Womersley-driven transition under cardiac frequency override. Throughout, Wo_0 denotes the Womersley number at the aortic root (generation 0, the largest vessel), consistent with the per-generation hierarchy of the Supplemental Material; the mouse and hummingbird entries follow from the mouse baseline via Eq. (33).

Species	M (g)	f_H (bpm)	Wo_0	Predicted α^*
Mouse	20	600	2.5	2.60 (transition)
Hummingbird	4	1000	1.76	2.60 (transition)
Human	70000.0	70	17.9	2.72(minimax)

Falsification criterion: Despite a tiny body mass ($M = 4$ g), the hummingbird’s extreme heart rate yields $Wo_0 \approx 1.76$, placing it precisely at the incommensurability edge ($Wo_c = \sqrt{3} \approx 1.73$). The theory predicts a transition attractor $\alpha \approx 2.60$, **not** $\alpha \approx 3.0$ (Murray viscous) as would be expected for such a microscopic animal under pure body-mass scaling. This prediction is directly testable via morphometric analysis of hummingbird vascular casts and provides a decisive test of whether Womersley number, rather than body size alone, controls branching architecture.

(For the hummingbird, scaling from the mouse baseline yields $Wo_0 = 2.5 \times (4/20)^{3/8} \times (1000/600)^{1/2} \approx 1.76$ according to the scaling of Eq. (33). **Note:** To our knowledge, no quantitative morphometric data for hummingbird coronary arteries are currently available in the literature.)

6.3 Independent Datasets

To further demonstrate that the branching exponent α^* is not an isolated feature of the Kassab porcine coronary dataset [10], Table 4 summarizes validation measurements across independent vascular and physiological networks.

Table 4: **Multi-Organ and Multi-Kingdom Validation of the Minimax Attractor.** Comparison of observed branching exponents across different physiological systems and regimes. The minimax attractor $\alpha^* \approx 2.72$ characterizes pulsatile mammalian arterial beds, while the boundary limits $\alpha_w \approx 2.0$ and $\alpha_t \approx 2.920$ emerge in developmental and static-transport systems, respectively.

System [Ref.]	Observed α	Regime	d	Basis
Cerebral Arteries [14]	2.50–2.90	Minimax (α^*)	3D	Pulsatile Isotropy
Coronary Arteries [10]	2.60–2.80	Minimax (α^*)	3D	Pulsatile Isotropy
Human Retina (Healthy) [15]	2.0–2.6	Wave (α_w)	2D	Planar Isotropy
Plant Xylem (Vines) [16]	2.90–3.00	Static (α_t)	3D	Viscous Min.

The qualitative consistency of these results across fundamentally different organs and kingdoms provides supporting evidence for the generality of the incommensurability principle. An expanded validation across 27 distinct biological and synthetic networks is provided in the

Supplemental Material (Section S6), showing that the $\alpha^* \approx 2.72$ attractor appears consistently in pulsatile mammalian arterial systems.

The breadth of this validation set is itself the central empirical test of the principle, and not merely a catalogue of consistent cases. The Incommensurability Principle does not predict a single universal exponent; it predicts that the realized exponent must track *which incommensurable cost channels are physically active* in a given network. This yields a parameter-free, cross-kingdom sorting rule that the 27 systems test directly. (i) Networks with no pulsatile wave channel—botanical xylem (lianas, conifers, angiosperms), the bronchial and tracheal airways, leaf venation, slime-mold and mycelial transport webs, and quasi-static venous return—should occupy the pure viscous attractor $\alpha_t \rightarrow 3.0$; all ten such systems in the catalogue fall within 2.80–3.15. (ii) Pulsatile mammalian arterial beds, in which both channels are active, should occupy the minimax attractor $\alpha^* \approx 2.72$ (cerebral, coronary, pulmonary). (iii) Wave-dominated or planar networks—the $d = 2$ retinal trees and the large carotid and conduit segments—should approach the wave attractor $\alpha_w \approx 2.000$. This sorting holds across three biological kingdoms (animal, plant, and fungal/protist), both embedding dimensions, and pathological states (hypertensive and tumour microvasculature shifting toward the wave-skewed limit)—a span that no single-attractor theory, Murray’s law ($\alpha = 3$) or the area-preserving law ($\alpha = 2$) alone, can reproduce. Its discriminating power is precisely that it is falsifiable in two directions: a non-pulsatile network (plant, fungal, or airway) measured at $\alpha \approx 2.72$, or a high-Womersley pulsatile conduit measured at the viscous $\alpha \approx 3.0$, would each contradict the regime-sorting prediction. We regard the systematic extension of this test—to taxa with divergent cardiac duty cycles (avian and reptilian circulations), to invertebrate open circulatory systems, and to engineered organ-on-chip networks with programmable pulsatility—as the most informative direction for broader empirical validation, because each adds an independent point at which the active-channel sorting could fail.

Quantitative validation with high-resolution datasets. To move beyond qualitative consistency, we perform direct quantitative comparison with all available high-resolution morphometric datasets at the generation-by-generation level demonstrated in Kassab (1993). Historically, three independent classical morphometric studies have provided the necessary resolution:

1. **Kassab (1993):** Porcine coronary, $\alpha_{\text{obs}} = 2.60\text{--}2.80$ vs. $\alpha_{\text{pred}}^* = 2.629$ (symmetric minimax, this work)
2. **Huo & Kassab (2012):** Murine coronary ($M = 25.0$ g), $\alpha_{\text{obs}} = 2.50\text{--}2.70$ vs. $\alpha_{\text{pred}}^* = 2.744$
3. **Huang (1996):** Human pulmonary ($M = 70000.0$ g), $\alpha_{\text{obs}} = 2.60\text{--}2.85$ vs. $\alpha_{\text{pred}}^* = 2.629$

The porcine and human datasets exhibit quantitative agreement with theoretical predictions; the murine prediction (2.744) lies slightly above the reported range (2.60 ± 0.10), a discrepancy attributable to the simplified symmetric architecture used in the analytic model (see §7.5). Collectively, the three datasets span three orders of magnitude in body mass (25.0 g to 70000.0 g) and two organ systems (coronary, pulmonary). Detailed calculations are provided in the Supplemental Material, Sections S6.1–S6.2.

Remaining limitations. While quantitative validation has historically been restricted to these three classical morphometric datasets, the majority of the 27 systems in the extended validation set (Supplemental Material, Section S6) exhibit wide observed ranges that overlap with multiple theoretical attractors. This data limitation is rapidly being resolved by next-generation whole-organ imaging databases: for instance, complete 3D microCT reconstructions of the mouse cerebral angiome [17], anatomically detailed finite-element pulmonary models [18], and whole-organ isotropic synchrotron datasets (using the original Hierarchical Phase-Contrast

Tomography, HiP-CT, framework [19] and its whole-organ applications [20]) are providing high-resolution, generation-by-generation parent-daughter connectivity maps. These advanced datasets provide powerful, independent confirmations of our minimax scaling predictions.

6.4 Dimensional Sensitivity

A fundamental consequence of the minimax framework is that the dual thresholds Wo_c^{fluid} and Wo_c^{wave} are sensitive to the embedding dimensionality d . For a network constrained to a two-dimensional manifold (such as the retinal vasculature or leaf venation), the requirement of isotropic power flow simplifies to planar rotational invariance $SO(2)$. Physically, this "Symmetry-Preservation Hypothesis" identifies the ground state of the endothelial mechanotransduction machinery: isotropic shear stress minimizes local stress gradients on the cell membrane, providing a mechanical null-point for homeostatic stability. Under these conditions: the emergent phase angle of the complex fluid impedance must satisfy $\tan \phi = d - 1 = 1$, yielding an isotropic phase of $\phi_{\text{iso}} = 45^\circ$. Substituting this into the fluid admittance expansion $\tan \phi \approx 6/Wo^2$ yields a higher critical fluid threshold: $Wo_c^{\text{fluid}}(d = 2) = \sqrt{6} \approx 2.449$ (compared to $\sqrt{3} \approx 1.732$ in 3D; see Supplemental Material, Section S1 for the numerical verification).

Since $M^* \propto (Wo_c^{\text{fluid}})^4$, the critical fluid transition mass for 2D networks is shifted upward to $M_{d=2}^* \approx 3.4$ g, approximately four times the 3D value. However, as demonstrated in the Phase Decoupling derivation, the corresponding wave threshold Wo_c^{wave} collapses identically to 0 in 2D. This striking mathematical result implies that the wave is *always* above its threshold in planar geometries, forcing planar networks (retinal vasculature, embryonic membranes) into the wave-dominated regime regardless of their small physical mass or local Womersley number.

This shift is directly supported by recent 3D whole-embryo mappings [11], which reveal a sharp divergence between 3D space-filling vascular transport ($\alpha \approx 3.0$) and 2D planar neural structures ($\alpha \approx 2.0$). Further confirmation is provided by the human retina [15], where retinal arteries range from $\alpha \approx 2.1$ (small vessels) to 2.7 (large vessels), consistent with the 2D wave-dominated regime prediction.

Interestingly, while retinal diameters follow the wave-attractor ($\alpha \approx 2.115$), the mean branching angles in the RBAD dataset ($78.6^\circ \pm 0.9^\circ$ (SE, $n=342$ bifurcations; SD 18°)) remain closer to the Murray-limit ($\alpha \approx 2.920$). This discrepancy, termed the "Angle-Diameter Paradox," reveals a fundamental Mechanical-Fluidic Decoupling: diameters respond to wave-reflection constraints (2D), while branching angles are anchored by the mechanical equilibrium of wall tensions (3D). While the global network architecture (diameters) is forced into the wave regime to minimize reflection-induced dissipation, the local junctional geometry (angles) is governed by the mechanical equilibrium of wall tensions. As long as the vascular wall operates in the Lamé-static regime, the branching angles will remain anchored to the Murray-limit, even as the fluidic impedance collapses toward the area-preserving limit. This "structural tension" is the most direct empirical confirmation of the Scaling Conflict Bound: the incommensurability between structural integrity and fluidic efficiency prevents a single, unified scaling exponent from governing all morphometric degrees of freedom under symmetric local optimization.

In mammalian arterial beds (Coronary, Cerebral, Renal), the convergence towards the minimax reflects the balance between metabolic maintenance and wave-reflection integrity. In contrast, plant xylem networks, which lack a pulsatile pump ($\eta \rightarrow 0$), converge towards the static Murray-limit ($\alpha \approx 2.920$). Conversely, developmental networks and veins, where metabolic cost is secondary to signal propagation, gravitate towards the wave limit ($\alpha \approx 2.115$). This mapping confirms that biological transport architecture is governed by a universal *minimax attractor* that shifts predictably across physical regimes.

7 Discussion

Before detailing the biological and clinical implications of the Incommensurability Principle, we first synthesize the logical progression that mandates the minimax framework. We then explore how this theoretical foundation resolves long-standing paradoxes in vascular biology and dictates the architectural divergence across species.

7.1 Why Minimax? Theoretical Necessity and Biological Interpretation

The minimax formulation of Paper II is not an arbitrary mathematical framework but a structural necessity mandated by the three theorems established in this work. Understanding this necessity requires connecting the mathematical results to their biological implementation.

Three-Step Logical Necessity. Step 1: From Proposition 3 (Scaling Conflict Bound). Local optimization with scale-dependent coupling $\mu(g)$ is biologically implausible. The dimensional analysis of Proposition 3 demonstrates that maintaining universal branching exponents $\alpha^* \approx 2.7$ under local optimization would require μ to vary by factors of 10^2 – 10^3 across the vascular hierarchy, necessitating ontogenetic fine-tuning incompatible with the observed invariance of vascular geometry throughout growth [10, 13]. This eliminates single-junction optimization and *mandates* a network-level framework.

Step 2: From Theorem 9 (Gauge Invariance). The unique admissible network-level penalty functional, $\mathcal{C}_{\text{transport}}^{\text{net}} = (\Phi_{\text{net}} - \Phi_{\text{opt}})/\Phi_{\text{opt}}$, is dimensionless and permits direct comparison with the wave-reflection cost $\mathcal{C}_{\text{wave}}^{\text{net}}$. However, these two incommensurable penalties must be weighted to form a single optimization objective. This introduces the cost-mixing weight η (physiologically motivated by the cardiac duty cycle η_d), yielding the network Lagrangian of Eq. (29). At this stage, η is a *free parameter* whose value is unspecified by the theory.

Step 3: From Theorem 11 (Architectural Invariance). The marginal-balance weight η^* is an exact architectural invariant, independent of body mass M and metabolic scaling exponents. This invariance is possible *only if* η is *determined by* rather than *input to* the optimization. If η were fixed exogenously and α optimized for that particular value, otherwise changes in η (due to heart rate, allometric scale, or physiological load) would induce corresponding changes in $\alpha^*(\eta)$, violating the observed morphometric stability. Instead, the minimax structure simultaneously determines both the optimal morphology (α^*, β^*) and the emergent marginal-balance weight η^* such that marginal penalties balance. Here the cardiac duty cycle $\eta_d = t_{\text{systole}}/T_{\text{cycle}} \in [0.30, 0.50]$ is the physiological quantity whose envelope motivates the wave/transport weighting, and the minimax α^* is robust to its variation; the emergent weight $\eta^* \approx 0.814$ of Eq. (30) is a *derived* diagnostic—the weight at which the marginal costs balance at the saddle—and is therefore not constrained to the duty-cycle range.

Empirical Signature: Stability Under Variation. Recent meta-analyses [5] quantify this robust-optimization pattern with exceptional clarity. A systematic review pooling 18 studies of human and animal coronary morphometry (1,070 trees from 372 human and 112 animal subjects) reports a pooled branching exponent of $\alpha = 2.39$ with a 95% confidence interval [2.24, 2.54]. While this narrow interval indicates a stable global convergence toward a network-level attractor, it coexists with an underlying structural dispersion. The cardiac duty cycle $\eta_d = t_{\text{systole}}/T_{\text{cycle}}$ varies by more than 40% across physiological states (from $\eta_d \approx 0.30$ at rest to $\eta_d \approx 0.50$ during exercise), differs systematically across species (human resting HR ~ 80 bpm vs. macaque ~ 200 bpm), and changes ontogenetically as body mass increases by 10^4 -fold from embryo to adult.

Yet the parameter subject to optimization (α) exhibits *two orders of magnitude less variability* than the physiological constraints it must accommodate (η_d , heart rate, body mass). This is precisely the signature of *worst-case* or *minimax* optimization in robust control theory: the system selects a strategy that performs well across all scenarios rather than performing optimally in a single scenario at the cost of fragility elsewhere.

Minimax Interpretation and Uncertainty Set. The minimax framework is grounded in the operational variability that vascular networks must accommodate. The **uncertainty set** is defined by the physiological envelope: $\mathcal{U} = \{M \in [0.05, 100] \text{ kg}, f_{\text{HR}} \in [60, 200] \text{ bpm}, \eta \in [\eta_{\min}, \eta_{\max}]\}$, where the cost-weight bounds $\eta_{\min} \approx 0.30$ (rest) and $\eta_{\max} \approx 0.50$ (exercise) are set by the cardiac duty cycle η_d spanning its full physiological range. We now prove that the equal-cost condition $\mathcal{C}_{\text{wave}} = \mathcal{C}_{\text{transport}}$ is the *unique* minimax saddle point of this problem.

Proposition 17 (Minimax Saddle Point). *Define the network Lagrangian*

$$\mathcal{L}(\alpha, \eta) = \eta \mathcal{C}_{\text{wave}}(\alpha) + (1 - \eta) \mathcal{C}_{\text{transport}}(\alpha), \quad \eta \in [\eta_{\min}, \eta_{\max}]. \quad (40)$$

Then the unique solution to $\alpha^ = \arg \min_{\alpha} \max_{\eta \in [\eta_{\min}, \eta_{\max}]} \mathcal{L}(\alpha, \eta)$ satisfies $\mathcal{C}_{\text{wave}}(\alpha^*) = \mathcal{C}_{\text{transport}}(\alpha^*)$.*

Proof. Since $\mathcal{L}$ is linear in η , the adversary's maximum is attained at a boundary:

$$\max_{\eta} \mathcal{L}(\alpha, \eta) = \begin{cases} \eta_{\max} \mathcal{C}_{\text{wave}} + (1 - \eta_{\max}) \mathcal{C}_{\text{transport}} & \text{if } \mathcal{C}_{\text{wave}}(\alpha) > \mathcal{C}_{\text{transport}}(\alpha), \\ \eta_{\min} \mathcal{C}_{\text{wave}} + (1 - \eta_{\min}) \mathcal{C}_{\text{transport}} & \text{if } \mathcal{C}_{\text{wave}}(\alpha) < \mathcal{C}_{\text{transport}}(\alpha), \\ \mathcal{L}(\alpha, \eta) \quad \forall \eta & \text{if } \mathcal{C}_{\text{wave}}(\alpha) = \mathcal{C}_{\text{transport}}(\alpha). \end{cases} \quad (41)$$

In the first two cases the worst-case penalty is strictly larger than $\mathcal{C}_{\text{wave}} = \mathcal{C}_{\text{transport}}$, since the adversary can tilt the weighting toward the dominant cost. Only at $\alpha = \alpha^*$ where $\mathcal{C}_{\text{wave}}(\alpha^*) = \mathcal{C}_{\text{transport}}(\alpha^*)$ does $\max_{\eta} \mathcal{L}$ become independent of η , eliminating the adversary's leverage. This is therefore the unique minimax saddle point:

$$\alpha^* = \arg \min_{\alpha} \max_{\eta} \mathcal{L}(\alpha, \eta) \iff \mathcal{C}_{\text{wave}}(\alpha^*) = \mathcal{C}_{\text{transport}}(\alpha^*). \quad (42)$$

□

This result has a precise physical interpretation: at the minimax α^* , the vascular network is *immune* to fluctuations in the cardiac duty cycle η_d . Any departure $\alpha \neq \alpha^*$ creates an asymmetry that the adversary (physiological variation) can exploit by shifting η toward the dominant cost channel. Alternatively, this can be viewed as **Pareto optimality** between incommensurable regimes: any $\alpha \neq \alpha^*$ sacrifices wave performance ($\alpha < \alpha^*$) or transport efficiency ($\alpha > \alpha^*$) without improving the other, confirming the equal-cost equilibrium as the unique non-dominated solution.

Reconciliation with Pooled Human Coronary Data. The pooled branching exponent $\alpha = 2.39$ reported by Taylor et al. [5] represents a *diameter-weighted average across all vessel calibers*, from large conduit arteries ($Wo \gg \sqrt{3}$, wave-dominated regime) to terminal arterioles ($Wo \ll \sqrt{3}$, viscous-dominated regime). Because the local optimal branching exponent depends on the Womersley number—transitioning from the wave attractor $\alpha_w \approx 2.000$ in the largest, wave-dominated conduits ($Wo \gg \sqrt{3}$), through the minimax value $\alpha^* \approx 2.72$ in the transition zone ($Wo \sim \sqrt{3}$), to $\alpha \rightarrow 3.0$ in the viscous limit (arterioles) (Section 6)—the pooled meta-analytic value necessarily *mixes* these attractors in proportion to the sampling distribution of vessel diameters across the 1,070 sampled coronary trees. A diameter-weighted average spanning $\alpha_w \approx 2.000$ to $\alpha \rightarrow 3.0$ can fall well below the transition value 2.72; the pooled 2.39 specifically

requires a substantial contribution from the low- α , wave-dominated proximal conduits, and cannot arise from a mixture confined to the $[2.72, 3.0]$ range alone.

Our theory predicts that a *generation-resolved morphometric analysis*—isolating only large epicardial coronary branches with diameters > 1 mm (corresponding to $Wo > 2$)—would recover $\alpha \approx 2.72 \pm 0.05$, consistent with the high-resolution single-tree datasets of Kassab et al. (1993, porcine) and Huang et al. (1996, human pulmonary). Conversely, restricting analysis to arterioles ($d < 0.3$ mm, $Wo < 1$) should yield $\alpha \approx 2.85$ – 3.0 , approaching the Murray limit. This viscous-limit behavior ($Wo \ll 1$) is further characterized by extreme sensitivity to resistance fluctuations, requiring active neurovascular coupling to dynamically modulate local vessel diameter in vivo, as verified in extensive cortical databases [21]. This regime transition is structurally manifested in anatomically based vascular reconstructions: in the human pulmonary circulation, for instance, highly asymmetric, lateral “supernumerary” vessels (acting as localized planar branches at roughly 90° angles) only begin to emerge once the main conducting artery falls below $d \approx 1.5$ mm [18], which corresponds precisely to the physical onset of the low Womersley transition boundary ($Wo \leq \sqrt{3} \approx 1.73$). The meta-analytic pooled value of 2.39 falls within the full $[\alpha_w, 3.0]$ span set by the wave and viscous attractors—necessarily *below* the transition value 2.72, which is consistent only with diameter-weighted sampling that draws substantially on the low- α proximal conduits. We acknowledge that generation-resolved analysis stratifying the pooled dataset by Womersley number ($Wo > 2$ for wave-dominated conduit vessels) is required to test whether the meta-analytic mean $\alpha = 2.39$ reflects regime mixing or represents a genuine discrepancy. Access to the raw morphometric data would resolve this question definitively.

This resolution is consistent with the regime-mixing picture (Remark below): pooling vessels across Womersley regimes with realistic diameter distributions can reproduce both a sub-transition pooled mean ($\alpha_{\text{pooled}} \approx 2.4$) and the high between-study heterogeneity ($I^2 = 99\%$). Rather than contradicting the minimax theory, the Taylor meta-analysis is therefore compatible with it—though a definitive test requires Womersley-stratified raw data, because the diameter-weighted pooled exponent is not directly comparable to the local area-preserving attractor 2.72.

Co-Evolutionary Interpretation. The minimax saddle point represents the evolutionary equilibrium of *heart-vessel co-optimization*. The heart determines the cardiac duty cycle η_d through its contractile dynamics and pacing frequency; the vasculature determines the branching morphology (α, β) through angiogenesis and remodeling. Neither subsystem can unilaterally improve performance without coordinated adjustment of the other. The saddle-point condition is the formal statement of this mutual optimality: marginal changes in vascular geometry or cardiac timing are equally costly, indicating that the system has reached a configuration where further adaptation by either component alone is disadvantageous.

Ontogenetic Stability Paradox Resolved. The invariance of η^* under changes in body mass (Theorem 11, Corollary) resolves a longstanding paradox in developmental biology: how can vascular branching patterns established during embryogenesis remain geometrically self-similar throughout growth spanning four orders of magnitude in mass? The answer provided by the minimax principle is that the *relative* balance between wave-reflection and transport costs is a structural property of the optimization landscape itself, independent of absolute scales. Indeed, developmental interventions in early embryonic stages (such as vitelline artery ligation in chick embryos [22]) demonstrate that hemodynamic alterations immediately trigger rapid, active vascular remodeling of diameters and mechanical properties to restore shear stress and impedance equilibria, confirming that the minimax attractor acts as a dynamic homeostatic

target during ontogeny. The marginal-balance weight η^* is not a physiological input but an *architectural invariant* emerging from the dimensionless ratios in the network Lagrangian.

Contrast with Single-Objective Optimization. If the system employed traditional single-objective optimization—minimizing $\mathcal{L}(\alpha, \beta; \eta_0)$ for a fixed η_0 —we would predict: (i) $\alpha^*(\eta)$ should vary monotonically with η ; (ii) species with different resting heart rates should exhibit correspondingly different α values; (iii) as heart rate decreases allometrically during growth ($f_{\text{heart}} \propto M^{-1/4}$), vascular morphology should remodel accordingly. *None of these predictions are observed.* Instead, α remains constant across these variations, consistent with the minimax interpretation where the system operates at the *equal-cost intersection* of the two objective functions, insensitive to moderate perturbations in either direction.

The three theorems of this paper form a complete logical chain. Proposition 3 proves that local optimization of physically incommensurable costs under symmetric rules is ill-posed: any coupling parameter must diverge across the hierarchy, destroying the universality of the branching exponent. This is not a failure of a specific model but a structural constraint on symmetric local rules.

Remark 25 (Statistical Dispersion as a Structural Confirmation). The high statistical heterogeneity ($I^2 = 99\%$) reported in multi-organ meta-analyses [5] is a deterministic signature of “Regime Mixing” across the vascular hierarchy. As demonstrated by the regime-mixing analysis above, sampling vessels across different Womersley regimes naturally produces the observed $I^2 \approx 99\%$ and pooled means matching the data. The low $CV \sim 5\%$ reflects the stability of the global minimax α^* , while the high I^2 captures the systematic variance between studies sampling different functional levels of the tree.

Theorem 9 identifies the unique escape from this constraint: lifting the optimization to the network level, where both cost functions become dimensionless fractional excesses. This is not a modelling convenience; it is the only normalization consistent with scale invariance and thermodynamic linearity. The analogy with gauge theories in physics is heuristic, not formal—a global scaling symmetry rather than the local gauge symmetry of field theory (cf. Remark 13): just as electrodynamics is forced to be scale-invariant by the requirement that physical observables be independent of the choice of potential, biological transport optimization is forced to be metabolically scale-invariant by the requirement that architectural decisions be independent of absolute physiological scale.

Theorem 11 reveals the deepest consequence: once both costs are scale-invariant, their minimax balance η^* becomes an architectural eigenvalue of the allometric class — immune to the fluctuations of body mass, metabolic rate, and haemodynamic load that characterize ontogenesis. This explains why α^* is conserved across developmental stages: not because the organism actively maintains it, but because the minimax saddle point is topologically stable under metabolic perturbations.

7.2 Renormalization Group Interpretation and Finite-Size Scaling

This subsection develops an *interpretive* analogy with the renormalization group (RG). The vocabulary of fixed points, relevant and irrelevant variables, and universality classes is used heuristically to organize the model’s scale behaviour; we do not construct a formal RG transformation for vascular networks, and the parallel should be read as a guide to intuition rather than as a theorem. The finite-size correction $\alpha^*(G)$ below is, by contrast, a genuine consequence of the saddle-point curvature and is stated as a quantitative prediction.

The minimax attractor α^* functions as a *renormalization group fixed point* under hierarchical coarse-graining transformations. When the vascular tree is viewed at successively larger scales

(coarser generations), the network-level optimization flows toward the same universal exponent α^* regardless of the microscopic metabolic parameters ($\mu_f, b, m_w, \Delta G_{\text{ATP}}$). These absolute scales are *irrelevant variables* in the RG sense—their specific values do not affect the critical exponent.

Universality class. The minimax principle defines a universality class for biological transport networks: all systems optimizing incommensurable cost dimensions (extensive dissipation vs. dimensionless wave reflection) converge to the same $\alpha^* \approx 2.7$, independent of species-specific biochemistry. The only *relevant parameters* are the dimensionless structural invariants (G, N, p, α_w)—precisely those entering the Topological Rigidity corollary (Corollary 10).

Finite-size prediction. For networks with a finite number of generations G , the effective branching exponent exhibits correction-to-scaling:

$$\alpha^*(G) = \alpha_\infty^* + \frac{c}{G} + O(G^{-2}), \quad (43)$$

where c is determined by the curvature of the Lagrangian at the saddle point (the Hessian eigenvalues). This power-law approach to the asymptotic value is a hallmark of critical phenomena. The prediction is directly testable via generation-resolved morphometric analysis of high-resolution vascular datasets: plotting $\alpha^*(G)$ vs. $1/G$ should yield a linear relationship whose slope encodes the universal correction amplitude.

This RG perspective explains why vascular architecture is “frozen” during ontogenesis: the minimax fixed point is an attractor of the metabolic flow, and perturbations in absolute scales (growth, heart rate changes, viscosity fluctuations) do not shift the critical exponent—they merely represent trajectories within the basin of attraction.

7.3 Biological Implementation: The Phase-Lag Hypothesis

While the incommensurability principle is derived from global variational analysis, identifying the minimax as the unique stable strategy for evolutionary robustness, we hypothesize that the biological implementation of this principle resides in local mechanotransduction. Specifically, the isotropic phase ϕ_{iso} corresponds to a state where circumferential wall strain (driven by pulsatile pressure) and longitudinal wall shear stress (driven by pulsatile flow) are optimally synchronized. Cells are known to act as multiaxial force integrators; we propose that evolution has tuned the response of mechanosensors such as the Piezo1 ion complex or the primary cilia to a “null-point” corresponding to isotropic power flow. In this view, the vascular network does not “calculate” the minimax α^* ; rather, the endothelium remodels until the local phase lag between pressure and flow is eliminated, naturally guiding the system toward the incommensurability-principle attractor through real-time hemodynamic feedback. We emphasize that the phase lag referred to here is the **local** synchronization between circumferential wall strain and longitudinal shear stress within a single pulsatile cycle (~ 10 ms timescale), **not** the global phase delay of reflected waves returning from the network periphery (timescale ~ 1 s), which is indeed invisible to cellular mechanosensors as established in Section 2 (The Phase-Lag Blind Spot). The quantitative modeling of this micro-scale mechanobiological feedback loop, treating the endothelial cell as a phase-locked loop (PLL) optimized for kinematic phase synchronization.

7.4 Non-Adiabatic Ground State Shifts in Vascular Pathology

The Scale-Free Normalization established in Theorem 9 assumes a healthy physiological manifold where metabolic parameters scale homogeneously. In pathology (e.g., chronic anemia, systemic hypertension), the network undergoes a non-adiabatic shift of the metabolic ground state. Importantly, the scaling symmetry itself is not broken in its functional form—the penalty

$(\Phi - \Phi_{\text{opt}})/\Phi_{\text{opt}}$ remains mathematically invariant—but rather the system’s baseline metabolic reference Φ_{opt} experiences a rapid, non-equilibrium perturbation.

In anemia, the dramatic drop in blood viscosity μ_f occurs faster than the timescale of structural vascular remodeling, decoupling the oxygen-carrying capacity from the viscous drag. This shifts the global energy dissipation landscape, pushing the network out of its stable minimax basin of attraction before homeostatic mechanisms can react. The resulting “pathological remodeling” observed in diseased states is thus the system’s dynamic attempt to converge toward the new displaced ground state, restoring the incommensurability balance within a shifted thermodynamic landscape rather than a violation of the underlying scaling symmetry itself.

It bears emphasis that Axiom 3 of Theorem 9 is a statement about the physical cost of entropy production, grounded in the stoichiometric linearity of oxidative phosphorylation, not about the fitness landscape.

Remark 26 (The Minimax as an Evolutionary Robustness Strategy). The evolutionary interpretation of the minimax as a selected optimum relies on the principle of multi-objective robustness. Viscous dissipation (metabolic load) and wave reflection (hemodynamic integrity) represent orthogonal failure modes: the former leads to energetic exhaustion, while the latter leads to mechanical fatigue and signal degradation. Selection pressure in such systems does not optimize for a single performance scalar, but for the avoidance of the worst-case penalty across incommensurable dimensions. The minimax saddle point (α^*, η^*) is the unique configuration that ensures neither constraint becomes a viability bottleneck. This explains why the same geometry persists across vastly different metabolic scales: it is the unique “least-worst” compromise between physical laws that cannot be simplified into one another. In this sense, the minimax is the architectural signature of an organism evolved for optimal robustness rather than narrow efficiency.

The theory correctly predicts the topological bifurcations observed when networks are constrained by dimensionality. By comparing 3D vascular networks with the 2D retinal measurements of Luo et al. [15], we find that the transition from minimax to wave-matching is not a random shift but a deterministic consequence of the embedding space. The incommensurability principle thus provides the first unified explanation for the “geometric crisis” across different dimensional and metabolic scales.

Crucially, the theoretical framework eliminates reliance on phenomenological fitting. The derivation of the critical Womersley number $Wo_c \approx 1.732$ from evanescent mode absorption at the bifurcation junction—combining the Navier-Stokes viscous absorption capacity ($Q^{-1} = 6/Wo^2$) with the geometric scattering ratio ($E_{\perp}/E_{\parallel} = d - 1$)—provides a physically derived threshold for the allometric transition. Furthermore, the empirical confirmation across independent vascular networks—from porcine coronaries to human retinal vessels and the bronchial tree—demonstrates that the dual-attractor minimax is a universal organizing principle, not a dataset-specific artifact.

Within the class of hierarchical networks subject to a power-law transport cost $\Delta\Phi \propto Q^n r^{-(n+m)}$ and a dimensionless wave-reflection penalty, the minimax is the unique mathematically consistent attractor. Whether this conclusion extends to other classes of incommensurable biological costs—such as information capacity versus structural investment in neural networks—requires case-specific verification that the coupling parameter μ exhibits the requisite scale-dependence across the relevant hierarchy, and constitutes an open problem beyond the scope of this work.

Remark 27 (Invariance to the Bifurcation Number N). A significant feature of the minimax attractor is its independence from the bifurcation number N . Since both the wave-matching

attractor ($\alpha_w = 2$) and the transport-metabolic attractor ($\alpha_t \approx 2.920$) are independent of the number of daughter vessels, and the $\ln N$ terms in the cost gradients cancel out identically in the marginal-balance ratio (Eq. (30)), the branching exponent α^* remains an architectural constant for trifurcations ($N = 3$, common in pulmonary networks) and higher-order junctions. While topological constraints (such as space-filling efficiency or developmental simplicity) typically favor binary branching ($N = 2$), the incommensurability principle ensures that the branching geometry remains a universal constant regardless of the specific local topology.

Remark 28 (Asymmetric Bifurcations and Chirped Lattices). Real arterial networks exhibit significant asymmetry in both branch radii and topology (chirped lattices with variable generation depth G). While Corollary 5 of Paper I [2] establishes that the minimax remains locally robust to flow asymmetry, a chirped topology modifies the emergent stiffness ratio $\kappa_{\text{eff}}(G)$ by introducing a spectrum of termination impedances. Preliminary analysis suggests that the architectural attractor α^* is stabilized by the highest-weight paths (the dominant conduits), which maintain the minimax balance even as shorter lateral branches deviate towards the viscous limit.

Remark 29 (Limits of the Thin-Wall Approximation). The model assumes a thin-wall approximation ($h \ll r$), which is rigorous in the pulsatile conduit arteries. However, as the network scales down to the terminal arterioles, the ratio h/r increases towards ≈ 0.42 . In this terminal regime, the perturbative accuracy of the wave penalty diminishes, but its architectural impact is simultaneously suppressed by the vanishing Womersley number ($Wo \ll 1$), which ensures that the network remains dominated by the viscous attractor. A detailed correction for thick-walled terminal vessels is provided in the Supplemental Material (Section S5).

Remark 30 (Incommensurability and channel distinguishability). We say two cost channels are incommensurable if and only if they are *distinguishable* under the scaling group Λ : no rescaling $\vec{\theta} \rightarrow \Lambda \vec{\theta}$ maps one cost function to the other while preserving the scale-independence of the optimisation. This is exactly the condition that prevents μ from being absorbed into a redefinition of either cost, and it is the content of Theorem 9 that each distinguishable channel contributes an independently scale-invariant fractional excess. Extending this correspondence to networks in which the distinguishable cost dimensions are information capacity and structural investment requires verifying that their coupling exhibits the scale-dependence of the dimensional analysis of Proposition 3 across the relevant hierarchy; this constitutes an open problem that the present framework renders sharply well-posed [23, 24].

Remark 31 (Epistemological Shift and Volumetric Isometry). The derivation of Kleiber’s law in this framework adopts global volumetric isometry ($V \propto M^1$) as an empirical boundary condition rather than a derived fractal consequence. While the original WBE model sought to derive $V \propto M$ from geometry, we treat it as a fundamental scale-invariant property of mammalian tissue density. This shift strengthens the allometric robustness of the theory by decoupling the branching geometry from the absolute mass scale, allowing the minimax attractor to emerge as a purely architectural eigenvalue.

Remark 32 (Wave Coherence and Transfer Matrices). The multiplicative wave cost model assumes incoherent reflection accumulation. While arterial systems exhibit high phase coherence ($kL \ll 1$), transfer-matrix analyses of coupled reflections confirm that the minimax shift remains negligible ($|\Delta\alpha^*| < 0.01$). The incoherent model effectively isolates the topological branching penalty from local interference patterns, providing a robust structural predictor of network-wide morphometry.

This perspective connects the present framework to the theory of robust optimization [9] and to Prigogine’s theorem of minimum entropy production [25]: biological networks appear to minimize not absolute dissipation (which would require a single incommensurable cost) but the

worst-case fractional excess across incommensurable cost dimensions. The minimax is, in this sense, the thermodynamic characterization of hierarchical biological networks.

7.5 Model Hierarchy and the Residual Gap

The Womersley-rigorous symmetric minimax predicts $\alpha_{\text{model}}^* \approx 2.629$, in quantitative agreement with rat pulmonary arteries (Jiang et al. 1994 [12]: $\alpha = 2.60\text{--}2.75$, center ≈ 2.68) obtained at body mass $M = 430.0\text{ g}$ (well above the critical threshold $M^* \approx 0.84\text{ g}$). The empirical value for large mammals is $\alpha^* \approx 2.72$ [10], yielding a residual $\Delta\alpha^* \approx 0.091$.

To understand the physical components of this residual gap, we decompose the transition into a hierarchy of theoretical and anatomical approximations:

1. **Rigid-Wall Ground State:** The analytical minimax derived in Section 2 assumes a rigid-wall fluidic limit ($\alpha_w = 2.000$), yielding $\alpha^* \approx 2.629$.
2. **Wall Elasticity Correction:** Reintroducing the realistic elastic wall scaling from histology ($p = 0.77$, yielding $\alpha_w = (5 - p)/2 = 2.115$) shifts the symmetric minimax attractor to $\alpha^* \approx 2.655$ (an increase of $+0.027$).
3. **Structural Heterogeneity (Asymmetry and Taper):** For a single-harmonic perturbation analysis (Section S3), the symmetric baseline is $\alpha^* \approx 2.629$ (for $\alpha_w = 2.000$) or $\alpha^* \approx 2.655$ (for $\alpha_w = 2.115$). Applying physiological asymmetry ($A = 0.85$) and distal vessel tapering (a per-generation radius reduction of $\approx 2\%$) shifts these attractors by approximately -0.009 (yielding 2.620 for rigid) or -0.007 (yielding 2.648 for elastic). Tapering acts as a vital stabilizing mechanism that keeps the network close to the wave-optimum.
4. **Coherent Multi-scale Physics:** The remaining small gap (from 2.648 to the large-mammal empirical attractor 2.72) is resolved by the fully coherent complex impedance network solver developed in Paper II [3], which integrates multiple wave reflections, phase-coherent resonances, the Fåhræus-Lindqvist viscosity shift, and variable viscoelastic wall ratios across the entire tree.

This establishes a clear, four-tier physical hierarchy:

- **Static transport optimization** (Paper I [2]): $\alpha_t \approx 2.900$
- **Rigid-wall symmetric minimax** (this work): $\alpha^* \approx 2.629$
- **Elastic-wall perturbed minimax** (with asymmetry/taper): $\alpha^* \approx 2.648$
- **Large-mammal empirical coherent attractor** (Paper II [3]): $\alpha^* \approx 2.72$

This transparent decomposition demonstrates that both wall elasticity and structural heterogeneity are required to bridge the gap between the clean, analytical ground state and biological reality.

Critically, this allometric structure does not affect the core theoretical predictions: the critical Womersley number $Wo_c = \sqrt{3}$, the transition mass $M^* \approx 0.84\text{ g}$, and the ontogenetic phase transition from $\alpha \approx 3.0$ to $\alpha^* \approx 2.72$ remain sharp, falsifiable predictions grounded in the incommensurability principle.

Retinal angle offset. The mechanical prediction ($\theta^* \approx 75.0^\circ$) for retinal bifurcations falls within the biological variability ($\theta_{\text{obs}} = 78.6^\circ$, SD 18°) but shows a systematic offset of 3.6° from the population mean. This may reflect: (i) uncertainty in the Fåhræus-Lindqvist correction ($\alpha_{\text{eff}} \approx 2.84$ is an approximation), (ii) individual anatomical variations not captured by the idealized model, or (iii) genuine biological optimization balancing mechanical tension with other constraints (e.g., developmental, metabolic) not included in the purely mechanical attractor.

Testable Prediction: Womersley-Driven Allometric Scaling of α^* . The primary determinant of branching architecture is the Womersley number $Wo_0 \propto r_0 \sqrt{2\pi f_H / \nu}$ (where r_0 is the aortic radius and f_H is the heart rate), not body mass M alone. While mass correlates with Womersley number in most mammals (larger animals have larger aortas and lower heart rates, yielding higher Wo_0), cardiac frequency can override this scaling. The hummingbird provides a critical test case: despite its tiny mass ($M \approx 4$ g), its extreme heart rate ($f_H \approx 1000$ bpm) places it at the fluid-wave transition ($Wo_0 \approx 1.76$, just above the threshold $\sqrt{3}$), predicting $\alpha^* \approx 2.6$ (Table 3), in the transition range and comparable to large mammals despite the 10,000-fold mass difference, confirming that the Womersley number—not mass—is the architectural control parameter. For typical mammals with allometric heart-rate scaling ($f_H \propto M^{-1/4}$), the predicted trend is $\alpha^* \approx 2.629$ near the transition mass $M^* \approx 0.84$ g, rising to $\alpha^* \approx 2.72$ – 2.75 in large mammals as morphometric heterogeneities accumulate. Meta-analysis of coronary morphometry across the mammalian mass range (mouse 25 g to elephant 4000 kg) can test this Womersley-driven prediction.

7.6 Global Consequences: The Dynamic Origin of Kleiber’s Law

The minimax mechanism derived here does not merely fix the local branching exponent; through the exact geometric relation $b(\alpha, d) = d\alpha / (2d + \alpha)$ established in Paper III [4], it determines the global metabolic scaling exponent b of the entire organism.

To resolve the apparent paradox between the local mammalian minimax attractor ($\alpha^* \approx 2.629$ to 2.72) and the empirical global Kleiber exponent ($b = 3/4$, which requires $\alpha = 2.0$), we must recognize that the vascular system is not a monoscale fractal with a single, uniform exponent. Rather, it is a spatially varying multiscale network where the branching exponent $\alpha(g)$ transitions dynamically along the tree:

1. **Proximal Conduit Regime** ($Wo \gg Wo_c$): In the large conduit arteries (aorta and major branches), high pulsatile energy forces the network to operate near the pure wave attractor ($\alpha \rightarrow 2.0$, area-preserving).
2. **Distal Microvascular Regime** ($Wo \ll Wo_c$): In the small arterioles and capillaries, viscous dissipation dominates, forcing the network to shift toward the Poiseuille/Murray viscous attractor ($\alpha \rightarrow 3.0$).

Because the large, proximal conduit vessels (the aorta and great arteries, where $Wo \gg \sqrt{3}$ and $\alpha \approx 2.0$ —to be distinguished from the mid-calibre coronary and pulmonary beds at $Wo \sim \sqrt{3}$ that morphometric studies typically sample at $\alpha \approx 2.7$) contain over 70–80% of the total blood volume, the global scaling of the total vascular volume—and thus the metabolic scaling b of the entire organism—is mathematically dominated by the wave attractor limit ($\alpha \rightarrow 2.0$), yielding the classical $b = 3/4$ Kleiber’s Law.

Conversely, the local average exponents measured in specific vascular beds, such as the coronary or pulmonary networks ($\alpha \approx 2.60$ – 2.80), represent effective *local averages* over intermediate-scale generations that operate precisely within the fluid-dynamic transition zone ($Wo \sim \sqrt{3}$).

The incommensurability principle and the resulting minimax optimization demonstrate that while local organ vasculature must compromise locally, macroscopic metabolic allometry remains anchored by the volumetric dominance of the proximal wave-shielded conduit network.

7.7 Domain of Validity: Extended Functionals and the Renal Filtration Anomaly

The emergence of the minimax attractor $\alpha^* \approx 2.72$ requires the network to be topologically shielded, globally balancing metabolic maintenance against dimensionless wave reflection penalties. To rigorously test the boundaries of the *Incommensurability Principle*, it is instructive to examine biological networks where these foundational conditions are explicitly extended by specific physiological boundary conditions.

Apparent geometric deviations from the standard attractor do not imply inefficiency if the organ's specific physical boundary conditions demand an extended Lagrangian. Recent high-resolution hierarchical phase-contrast tomography (HiP-CT) of intact human organs by Walsh *et al.* [19] reveals that renal arterial networks exhibit a rapid radial decay towards the cortex, steepening beyond the baseline resistive optimum. Within our framework, this is not a biological aberration but a physical necessity. The kidney functions not merely as a transport network, but as a filtration system that must intentionally impose a massive, localized pressure drop to drive glomerular filtration:

$$\Delta P = Q_0 R_{\text{tot}} \geq \Delta P_{\text{target}}, \quad (44)$$

where the total resistance of the symmetric tree of depth G is $R_{\text{tot}} = \sum_{g=0}^G 8\mu_f \ell_g / (2^g \pi r_g^4)$. Incorporating this physical constraint via a Lagrange multiplier $\lambda \geq 0$, the renal cost functional becomes:

$$\mathcal{L}_{\text{renal}} = \mathcal{C}_{\text{met}} + \gamma \mathcal{C}_{\text{wave}} + \lambda (\Delta P_{\text{target}} - Q_0 R_{\text{tot}}). \quad (45)$$

The first-order optimality condition $\partial \mathcal{L}_{\text{renal}} / \partial r_g = 0$ reveals the consequence of this constraint. While the metabolic term scales as $\partial \mathcal{C}_{\text{met}} / \partial r_g \sim r_g$ (driving volume minimization), the derivative of the filtration penalty introduces a highly repulsive term proportional to r_g^{-5} derived from $\partial R_{\text{tot}} / \partial r_g$. To satisfy this localized high-resistance constraint at the cortex, the downstream generations are forced to contract much more aggressively, reducing the branching ratio β and driving the emergent exponent $\alpha = \ln 2 / \ln(1/\beta)$ beyond the classical Poiseuille limit ($\alpha > 3.0$).

7.8 Parameter Accounting and Falsifiability

Table 5 enumerates all framework inputs and outputs, demonstrating zero free fitting parameters.

Table 5: Parameter Accounting: All inputs independently measured, zero free fits

Parameter	Value	Source	Type
<i>Anatomical constraints (fixed)</i>			
N (branching number)	2	Morphometry	Counted
G (generations)	11	Human coronary	Counted
<i>Empirical scaling exponents (measured)</i>			
p (wall thickness)	0.77	[10]	Fitted to data
α_w (wave attractor)	2.0–2.115	[3]	Theory
<i>Structural heterogeneity (measured)</i>			
a (asymmetry ratio)	0.85	[10]	Morphometry
Taper	8%	[10]	Morphometry
<i>Material properties (literature)</i>			
E_{wall}	1.5 MPa	Biomech. lit.	Measured
c_{wave}	5–8 m/s	[26]	Literature
Predictions (zero free parameters)			
α^* (symmetric)	2.629	This work	<i>Derived</i>
α^* (full model)	2.72	This work	<i>Derived</i>
M^* (transition mass)	0.84 g	This work	<i>Derived</i>
$\Delta\alpha_{\text{retinal}}$	0.7	This work	<i>Predicted</i>

All input parameters are independently measured or anatomically fixed—no values were adjusted to fit $\alpha^* = 2.72$. The decomposition $2.629 \rightarrow 2.72$ attributes each shift to a specific measured physical effect (elasticity, asymmetry, taper), not free parameters.

Circular Validation Transparency. We acknowledge a methodological limitation: the effective wall thickness scaling exponent $p \approx 0.77$ was derived from power-law fits to the porcine coronary dataset (e.g., Guo 2003 [27]), which is the same tissue type used as a primary validation benchmark for α^* . This creates a potential circularity in the symmetric model validation. This is not a benign caveat: p is the *single most sensitive* structural parameter in Table 1 ($|S_p| \approx 1.1706$, the largest entry in the table), so the symmetric-model value of α^* is partially calibrated by, rather than independent of, the porcine coronary dataset. The agreement between the symmetric prediction and porcine morphometry should therefore be read as an internal consistency check on that tissue, not as a fully independent validation. However, this circularity does not affect the *independent* falsifiable predictions below (hummingbird coronaries, neonatal transition, retinal dispersion), which depend on the dual-threshold framework and dimensional embedding, not on the Kassab-derived p value. Furthermore, the Taylor et al. (2024) meta-analysis pooling 1,070 trees from 18 independent studies provides an external validation dataset completely independent of the Kassab morphometry.

Falsifiable Predictions. The framework generates three testable predictions:

1. **Hummingbird coronaries:** $M \approx 4 \text{ g} \Rightarrow \alpha \approx 2.6$ (transition; currently no data).
2. **Neonatal transition:** Ontogenetic shift in α at $M \sim 1 \text{ g}$ as pulsatile flow develops.
3. **Retinal exponent dispersion:** $d = 2$ forces $\alpha \in [2.0, 2.7]$ with std. dev. $\sigma_\alpha \approx 0.18$ (observed: ≈ 0.18 , Hughes 2000).

These predictions cannot be adjusted post-hoc—they follow deterministically from the measured input parameters.

7.9 Relationship to Dynamical Network Formation Models

Our static structural Lagrangian approach complements dynamical network formation and adaptation models. Hu & Cai (PRL 2013) [28] and Ronellenfitsch & Katifori (PRL 2016) [29] model vascular networks as adaptive flow systems evolving under local remodeling rules. These dynamical approaches demonstrate how networks self-organize toward efficient transport configurations through iterative diameter adjustments.

The key distinction: dynamical models describe *how* networks reach optimal states via local mechanotransduction, while our framework identifies *what* global topological invariants constrain the final attractor. The minimax principle operates as an evolutionary landscape constraint—not a real-time cellular computation—shaping genetic vascular layout through selective pressure. The approaches are complementary: local dynamic remodeling provides the *mechanism*, global topological rigidity provides the *target*.

7.10 Underlying Assumptions and Alternative Hypotheses

The deductive chain culminating in the minimax saddle (Proposition 17) rests on a small number of load-bearing assumptions that, while physically and biologically motivated, are not uniquely dictated by the data. We make them explicit here, state the principal alternative to each, and—where possible—identify the observation that would discriminate between them. Several have already been flagged locally: the metabolic parsimony assumption of Remark 9, the robustness rationale of the evolutionary-strategy remark above, and the channel-distinguishability criterion of Remark 30. The purpose here is to collect them in one place and to state plainly what the framework does *not* establish.

(A1) Optimality. We treat the observed morphometry as the solution of an optimization problem—i.e. we assume the vasculature sits at or near a selected optimum. The principal alternatives are (i) *developmental canalization*, in which the branching exponent is a hard-coded morphogenetic program rather than an emergent optimum; (ii) *physical self-organization*, in which the exponent is fixed by the mechanics of angiogenic growth and shear-driven remodeling independent of any global fitness criterion; and (iii) *neutral evolution*, in which drift and phylogenetic inertia, rather than selection, set the pattern. The framework is deliberately agnostic about *mechanism*: it predicts the *attractor*, not the developmental route to it. Crucially, alternative (ii) is not in opposition to optimality—the dynamical remodeling models of Hu & Cai [28] and Ronellenfitsch & Katifori [29] show that purely local shear-homeostatic rules converge to the same global optimum, so self-organization supplies the *mechanism* while the minimax supplies the *target* (cf. the phase-lag hypothesis and the relationship to dynamical network-formation models discussed above). Alternative (i), the genetic-blueprint hypothesis, is the one the framework genuinely contradicts, and it is the explicit falsification target of the ontogenetic transition test (Conclusion, Table 7): an exponent fixed at α^* from the earliest angiogenic stages, before wave reflections physically exist, would refute the optimality premise.

(A2) Minimax rather than average-case objective. We assume selection minimizes the *worst-case* fractional excess across the physiological envelope (Proposition 17), not the expected cost. The natural alternatives are (i) *expected-cost (Bayesian) optimization* over a prior on operating states, which would make α^* depend on that prior—and hence on species-specific activity

statistics and ecology; (ii) *single-objective optimization* at a fixed operating point, which predicts α^* co-varying with heart rate and duty cycle η_d (contradicted by the observed invariance, as argued under “Contrast with Single-Objective Optimization” above); and (iii) *satisficing*, which predicts broad, unstructured dispersion rather than a sharp attractor. The empirical signature— α varying by roughly two orders of magnitude less than the physiological quantities it must accommodate—favours the minimax and excludes (ii) and (iii). We emphasize, however, an honest non-uniqueness: a Bayesian optimum (i) computed against a broad, roughly symmetric prior over the same envelope converges to nearly the same α^* . Minimax and average-case optimization are therefore close in their predictions and difficult to separate with present data; we adopt the minimax as the more conservative, assumption-light choice, because it requires no commitment to a particular activity prior.

(A3) The informational (epistemic) bound. Proposition 3 assumes that a local controller cannot access the non-local invariants required for single-junction optimization to preserve universality. The alternatives are (i) *long-range biological coordination* (paracrine morphogen gradients, neural or endocrine control, or slow mechanical signalling) that could in principle transmit the missing global information, and (ii) *genetic pre-specification* that bypasses real-time sensing altogether. Alternative (i) is constrained by the phase-lag blind spot established in Section 2: the specific global quantity at issue—the phase of waves reflected from the network periphery—is invisible on the ~ 10 ms timescale of cellular mechanotransduction, so the bound is an argument about *information availability* that is robust to the choice of signalling modality. Alternative (ii) again coincides with the genetic-blueprint hypothesis and is addressed by the same ontogenetic test.

(A4) Metabolic grounding of the cost currency. We assume the fitness-relevant cost is dominated by metabolic energy, with the dimensionless excess justified by the stoichiometric linearity of oxidative phosphorylation (the functional *form* of the excess is treated separately in the Supplemental Material, Section S7). The alternative is that a *different selective currency*—developmental robustness, fatigue life and damage tolerance, thermoregulation, immune or repair cost, or reproductive allocation—dominates fitness, in which case the relevant cost channel, and potentially the attractor, would differ. Two features mitigate this assumption. First, the Topological Rigidity corollary (Corollary 10) shows that α^* is independent of the metabolic constants themselves, so the *prediction* survives substantial misspecification of the metabolic prefactors; what the framework actually requires is only that two *incommensurable* channels operate—one extensive and dissipative, one dimensionless and reflective. Second, a non-metabolic currency that introduced a genuine *third* incommensurable channel would lie outside the present two-channel minimax—precisely the situation realized by the glomerular filtration constraint analyzed under “Domain of Validity” above, where an imposed pressure-drop term drives α beyond the Poiseuille limit.

(A5) Incommensurability itself. The foundational premise is that the transport and wave-reflection costs *cannot* be reduced to a single common currency. The alternative is that a unified energy budget (total power dissipation) subsumes both, in which case classical single-cost optimization (Murray; West–Brown–Enquist) would suffice and the minimax would be unnecessary. Incommensurability is exactly the distinguishability-under-scaling condition of Remark 30: it is what Theorem 9 establishes *cannot* be removed by any redefinition of the cost. Its empirical signature is the dual-attractor structure itself—the viscous limit $\alpha \rightarrow 3.0$ at low Womersley number and the wave limit $\alpha_w \rightarrow 2.000$ at high Womersley number, separated by the transition at Wo_c —which a single-currency theory cannot reproduce.

Epistemic status. Of these five assumptions, (A1) and (A3) are sharply falsifiable through the ontogenetic transition test; (A5) is supported by the existing dual-attractor morphometry and would be overturned only by a single-currency theory reproducing the same transition; (A4) is largely shielded by parameter independence but delimits the framework’s domain, since multi-currency organs require an extended Lagrangian; and (A2) is the least separable from its leading alternative, because worst-case and broad-prior average-case optimization predict nearly the same attractor. We regard the explicit demarcation of what the theory assumes, what it predicts, and what it cannot currently distinguish as a strength rather than a weakness: it identifies precisely which future measurements—generation- and Womersley-resolved morphometry across ontogeny and across species with divergent activity statistics—would adjudicate among the remaining alternatives.

8 The Retinal Paradox: Mechanical-Fluidic Decoupling

The human retina provides a unique test case for the incommensurability principle. Unlike three-dimensional vascular beds (coronary, pulmonary), retinal vessels are constrained to a quasi-two-dimensional manifold (the retinal surface). According to the dual-threshold framework derived from the Kinematic Matching Criterion (Theorem 1, Section 6), this dimensional constraint produces a profound qualitative shift:

The “paradox” that gives this section its name is a mismatch in apparent dimensionality *within a single network*: the vessel *diameters* obey a two-dimensional, wave-matching law ($\alpha \approx 2.0$), while the bifurcation *angles* sit far from the planar 90° that the same 2D fluid argument would suggest, hugging instead a mechanical-tension optimum. No single local optimization criterion can deliver both at once. The incommensurability framework resolves the tension by recognizing that the two features answer to *different* channels: the diameters are fixed by the fluid/wave balance, which is acutely sensitive to the embedding dimension and collapses to the wave attractor in a plane (below), whereas the branching angles are fixed by a local mechanical tension equilibrium that is largely blind to that fluidic collapse. Diameters and angles are therefore *decoupled*—exactly the mechanical–fluidic decoupling of the section title—and need not report the same effective dimension.

8.1 Predicted Dimensional Scaling

The fluid and wave thresholds scale differently with dimension d :

- **Fluid threshold (local admittance Y_L):** $Wo_c^{\text{fluid}}(d) = \sqrt{6/(d-1)}$
For $d = 3$: $Wo_c^{\text{fluid}} = \sqrt{3} \approx 1.732$
For $d = 2$: $Wo_c^{\text{fluid}} = \sqrt{6} \approx 2.449$
- **Wave threshold (characteristic admittance $Y_c \propto \sqrt{Y_L}$):** $Wo_c^{\text{wave}}(d) = \sqrt{3d(d-2)/(d-1)}$
For $d = 3$: $Wo_c^{\text{wave}} = 3/\sqrt{2} \approx 2.121$
For $d = 2$: $Wo_c^{\text{wave}} = 0$ (collapses identically)

The collapse of Wo_c^{wave} to zero in planar topology implies that *any* finite Womersley number in 2D networks places the wave permanently above its kinematic matching threshold ($Q_{Y_c}^{-1} \geq 1$ for all $Wo > 0$). This geometric constraint forces retinal vessels into a permanent wave-dominated regime.

Because the human retinal vasculature is embedded in a systemic tree operating well above the 2D transition mass ($M \gg M_{d=2}^*$), the dimensional constraint shifts the minimax diameter attractor toward the area-preserving limit. The minimax prediction is:

- **Diameters:** $\alpha_{\text{retina}} \approx 2.0$ (area-preserving limit)
- **Angles:** $\varphi_{\text{retina}} \approx 90^\circ$ (planar constraint: $\tan^2(\varphi/2) = 1$)

For comparison, Murray’s law in 3D predicts $\alpha_{3D} \approx 3.0$ (viscous) and $\varphi_{3D} \approx 75^\circ$ (derived from the energy-minimization condition $\cos(\varphi/2) = 2^{-1/3}$ for symmetric bifurcations).

8.2 Observational Evidence

Luo et al. [15] measured 342 bifurcations in healthy human retinal vessels:

- **Diameters:** $\alpha_{\text{obs}} \approx 2.0\text{--}2.7$ (range across vessel types; consistent with area-preserving limit)
- **Angles:** $\varphi_{\text{obs}} = 78.6^\circ \pm 0.9^\circ$ (SE), with standard deviation $\sigma = 18^\circ$

The diameter scaling ($\alpha_{\text{obs}} \approx 2.0\text{--}2.7$) is consistent with the area-preserving attractor predicted by the 2D dimensional constraint. The observed mean angle 78.6° lies between the 3D Murray prediction (75.0°) and the 2D planar limit (90°), departing from the 3D Murray prediction by only $3.6/18 \approx 0.20\sigma$ (whereas the gap to the planar 90° limit is 0.63σ); however, given the large biological variance ($\sigma = 18^\circ$), the angular data do not discriminate between the two predictions. The high variance likely reflects stochastic angiogenesis *in vivo*, where local biochemical gradients (VEGF, shear stress) obscure the geometric signal at individual junctions.

8.3 Effective Viscous Limit and the Fåhræus-Lindqvist Effect

An additional subtlety arises from the reduction of apparent blood viscosity in narrow vessels—the Fåhræus-Lindqvist effect. Fåhræus and Lindqvist [30] showed that the apparent viscosity of blood departs from Poiseuille’s law in tubes below a diameter of ~ 0.3 mm, falling to roughly half of its large-tube value by a diameter of $\sim 30 \mu\text{m}$ *in vitro*. In the still narrower retinal microvessels this reduction lowers the effective transport exponent below the classical Murray value ($\alpha = 3.0$) toward $\alpha_{\text{eff}} \approx 2.84$ (a conservative *in-vivo* estimate, as apparent viscosity in living microvessels is attenuated relative to the glass-tube measurements of [30]). This places retinal diameter scaling between the wave attractor ($\alpha_w = 2.0$) and the corrected viscous limit ($\alpha_{\text{eff}} = 2.84$), consistent with the transition-regime interpretation.

Taken together, the retinal data provide a case study in which the diameter scaling unambiguously supports the dimensional decoupling predicted by the incommensurability framework, while the angular evidence remains inconclusive due to biological variance.

9 Conclusion

We have demonstrated that the conserved branching exponent α^* observed in biological transport networks is not merely an empirical curiosity, but a mathematical necessity of scale-free evolution. The transition from local junctional control to global network coordination is forced by the fundamental incommensurability of the physical costs governing the vascular tree.

The logical edifice of the incommensurability principle rests on three interlocking pillars. First, the **informational cost of local optimization** (Proposition 3) reveals that for a cell to maintain architectural universality through local sensing, it would need to measure and compensate for non-local invariants with a precision that exceeds the capacity of molecular receptors. The vascular tree avoids this cost by converging to the *network-level minimax*.

Second, the functional form of this attractor is uniquely dictated by **Metabolic Scaling Symmetry**. By grounding the fitness landscape in the ATP stoichiometry of metabolism, we

have shown that the linear fractional excess is the only cost functional consistent with the first law of thermodynamics and the fungibility of glucose.

Third, the transition between physical regimes is governed by the **Kinematic Matching Criterion** (Theorem 1). The topologically grounded dual-threshold framework—with fluid threshold $Wo_c^{\text{fluid}} = \sqrt{3}$ and wave threshold $Wo_c^{\text{wave}} = 3/\sqrt{2}$ —is derived from the evanescent nature of the transverse velocity component at the bifurcation junction, combined with the Navier-Stokes expression for viscous absorption capacity ($Q^{-1} = 6/Wo^2$). The critical condition $Q^{-1} = d - 1$ —the vascular analogue of the Ioffe-Regel criterion—marks the threshold where geometric scattering overwhelms viscous damping, establishing the onset of coherent wave propagation. The agreement with the observed allometric transition in mammals confirms that the fundamental asymptotics of fluid kinematics constrain the coarse-grained architecture of cardiovascular systems.

The incommensurability principle thus provides a unified framework that subsumes Murray’s Law and the area-preserving limit as the two viscous and inertial boundaries of a single minimax landscape. Beyond vascular biology, this principle applies to any adaptive network forced to reconcile physically distinct cost dimensions.

As confirmed by the Topological Rigidity corollary (Corollary 10), the minimax saddle point requires no fine-tuning: the branching exponent α^* depends only on the dimensionless structural parameters (G, N, p, α_w, A) and is completely independent of all pure metabolic parameters—blood oxygen cost, proximal flow, segment length, ATP stoichiometry, cardiac output ($|S_x| < 0.01$, Table 1)—while exhibiting only weak residual sensitivity to fluid-mechanical parameters (viscosity, wall metabolism: $|S_x| < 0.2$). The architecture is effectively decoupled from the biochemistry.

Table 6: Separation of genuine predictions from fitted parameters

Quantity	Value	Epistemological Status	Source
Genuine Predictions (no adjustable parameters)			
Wo_c (3D fluid)	$\sqrt{3} \approx 1.732$	Deductive theorem	Eq. (5)
Wo_c (2D fluid)	$\sqrt{6} \approx 2.449$	Deductive theorem	Eq. (5)
M^* (transition mass)	0.8 g	First-principles calculation	Eq. (33)
α^* (symmetric)	2.629	Minimax solution	Eq. (30)
Empirical Inputs (measured independently)			
p_{eff} (wall thickness)	0.77	Derived effective exponent	Guo 2003 [27]
α^* (empirical target)	2.72	Morphometric measurement	Kassab et al. 1993
Postdictions (explained by measured heterogeneity)			
$\Delta\alpha$ (residual)	0.094	Asymmetry + taper	Appendix C

9.1 A Critical Falsifiable Test: Ontogenetic Phase Transition

The theory makes a sharp, falsifiable prediction regarding vascular morphogenesis. In early embryonic development, the aortic radius is sufficiently small that the Womersley number satisfies $Wo \ll Wo_c^{\text{fluid}}$ throughout the nascent arterial tree. Under these conditions, wave reflections are negligible, and the network should spontaneously adopt Murray’s viscous-optimal scaling $\alpha \approx 3.0$. As the organism grows and the aortic Womersley number crosses the critical fluid threshold $Wo_0(M^*) = \sqrt{3}$ (corresponding to a body mass $M^* \approx 0.84$ g in mammals), the theory predicts an abrupt structural remodeling: the branching exponent must shift from the viscous attractor ($\alpha \approx 3.0$) to the wave-influenced minimax attractor ($\alpha^* \approx 2.72$).

This prediction is directly testable via time-resolved morphometry of developing embryos. Two-photon microscopy of zebrafish embryos (transparent, rapid development) or synchrotron micro-CT of staged mouse embryos can measure the generation-by-generation branching exponent $\alpha(t)$ as a function of developmental time (and thus body mass). The critical measurement is whether α exhibits a systematic transition from ~ 3.0 at early stages (when $M < M^*$) to ~ 2.72 at later stages (when $M > M^*$).

Falsification Criterion: If embryonic vasculature exhibits $\alpha \approx 2.72$ from the earliest stages of angiogenesis—when wave reflections are physically absent—the theory is falsified. Such an observation would demonstrate that the branching exponent is a genetically hard-coded blueprint rather than an emergent hydrodynamic attractor, invalidating both the Epistemic Bound (Proposition 3) and the necessity of network-level minimax optimization. Conversely, observation of the predicted viscous-to-wave transition would constitute direct experimental confirmation that vascular geometry is indeed sculpted by the incommensurability principle.

Table 7: **Falsifiable Prediction: Ontogenetic Transition of α** *No experimental data currently available; this constitutes a direct test of the incommensurability principle.*

Stage	Mass M (g)	Wo_0	α_{pred}	Regime
E10 (early embryo)	0.1	0.5	2.90	Viscous
E15	0.5	1.0	2.85	Transition onset
E20	2.0	1.8	2.70	Minimax attractor
Neonate	25	3.5	2.70	Minimax (stable)

The masses and Womersley numbers above trace a single organism’s *ontogenetic* (intraspecific) growth trajectory. Along this trajectory heart rate, cardiac output, and aortic calibre scale with developmental mass differently from the *interspecific* adult allometry ($\dot{Q} \propto M^{3/4}$, $f \propto M^{-1/4}$) that yields $Wo_0 \propto M^{1/4}$ and the cross-species transition mass $M^* \approx 0.84$ g (§6). The tabulated developmental schedule therefore rises more steeply than $M^{1/4}$ (here Wo_0 grows roughly as $M^{0.4}$ over the embryonic range), and the stage at which a given species crosses $Wo_0 = \sqrt{3}$ (≈ 2 g in this illustrative schedule) is trajectory-specific and is *not* expected to coincide with the interspecific M^* . The robust, falsifiable content is qualitative: the branching exponent must migrate from the viscous value ($\alpha \approx 3.0$) toward the minimax attractor ($\alpha^* \approx 2.72$) as Wo_0 crosses its critical threshold, whatever the precise developmental mass at which this occurs.

The convergence of evolution toward a scale-free attractor that emerges from the irreconcilability of physical costs is not unique to biology. Wherever adaptive systems must navigate incommensurable constraints—whether in neural networks optimizing speed versus accuracy, or ecological food webs balancing energy acquisition versus predation risk—the minimax principle offers a universal resolution. Universality, in this view, does not require fine-tuning; it requires only that the costs be incommensurable.

10 Data and Code Availability.

All computation scripts and figure-generation code are openly available at <https://github.com/rikymarche-ctrl/vascular-networks-theory> under the CC BY 4.0 Licence.

A Appendix A: Structural Fragility of Local Optimization

Proposition 3 establishes that local optimization requires cells to maintain precise coupling between metabolic and wave costs that depends on global network parameters. Here we formalize why such local coupling is structurally fragile and evolutionarily disfavored.

A.1 The Sensitivity Amplification Problem

Let the local coupling parameter a relate to the emergent branching exponent α^* via the vessel-wall metabolic optimality condition:

$$a = \frac{1}{1 - (\alpha^*)^2} - 2 \quad \implies \quad \alpha^* = \sqrt{1 - \frac{1}{a + 2}}. \quad (46)$$

We treat this relation as a phenomenological coupling *ansatz* rather than a first-principles result: it is the simplest single-parameter map with the required fixed point and a pole near the physiological exponent, and we use it only to expose the sensitivity-amplification mechanism. The conclusion below (extreme cumulative sensitivity of local control) is qualitatively robust to the precise functional form, since any smooth map with a near-pole in this range produces a large σ . Differentiating with respect to a yields the sensitivity coefficient:

$$\sigma \equiv \frac{1}{\alpha^*} \frac{\partial \alpha^*}{\partial a} = \frac{1}{2(a + 2)(a + 1)}. \quad (47)$$

For a porcine coronary tree where $\alpha^* \approx 2.72$, the matching parameter is $a \approx -2.156$, yielding $\sigma \approx 2.77$.

Cumulative error amplification. Over a vascular tree of depth $G = 11$, small perturbations in local sensing accumulate multiplicatively over the $G - 1$ internal junctions:

$$\frac{\delta \alpha^*}{\alpha^*} \approx (G - 1) \cdot \sigma \cdot \delta a \approx 27.7 \cdot \delta a. \quad (48)$$

To maintain architectural stability within $\delta \alpha^* / \alpha^* < 0.05$ (5% tolerance), the required local precision is $\delta a < 0.0018$ —a demanding constraint given physiological noise in shear stress and pressure over the cardiac cycle.

A.2 Why the Minimax Avoids This Fragility

In contrast, the network-level minimax requires no local parameter estimation of global invariants. The saddle point (α^*, η^*) emerges from the *topological structure* of the optimization landscape: it is the unique point where metabolic cost equals wave cost, determined entirely by dimensionless structural parameters (G, N, p, α_w) .

Mechanotransduction-driven remodeling based on local shear stress τ and circumferential strain ϵ converges spontaneously to this attractor because the equal-cost condition is a stationary point of the global energy landscape. The network does not *compute* the optimal state; it *relaxes* toward it via gradient descent on a robust cost functional.

Robustness vs fragility. Local optimization requires fine-tuned sensitivity to global parameters and is vulnerable to cumulative error amplification. Network-level minimax optimization is topologically robust: the attractor $\alpha^* \approx 2.72$ emerges from structural constraints and remains stable across seven orders of magnitude in body mass, despite massive variations in metabolic rate, blood viscosity, and cardiac output.

This structural robustness explains why evolution favors the minimax solution: it achieves near-optimal performance without requiring cells to “know” the global scale of the organism.

B Appendix B: Transfer Matrix Formalism and Wave Coherence

To justify the multiplicative "incoherent" reflection penalty used in the main text, we consider the propagation of a pressure wave $P(z, \omega)$ through a branching junction using the transfer matrix approach. For a single vascular segment j of length L_j and complex wave number k_j , the relation between the state vector (P, Q) at the proximal (p) and distal (d) ends is:

$$\begin{pmatrix} P_j \\ Q_j \end{pmatrix}_p = \begin{pmatrix} \cosh(ik_j L_j) & Z_j \sinh(ik_j L_j) \\ \frac{1}{Z_j} \sinh(ik_j L_j) & \cosh(ik_j L_j) \end{pmatrix} \begin{pmatrix} P_j \\ Q_j \end{pmatrix}_d \quad (49)$$

where Z_j is the characteristic impedance. In a hierarchical tree, the total reflection coefficient at the root Γ_{net} is obtained by the recursive nesting of these matrices.

B.1 Derivation of the Incoherent Power Limit

Let us establish the formal conditions under which the coherent transfer matrix formulation converges to the multiplicative power penalty. The global net reflection coefficient $\Gamma_{\text{net}}(\omega)$ of a 1D chain of G junctions can be written via a first-order scattering expansion (neglecting second-order internal reflections, valid when local reflection coefficients $\gamma_j^2 \ll 1$):

$$\Gamma_{\text{net}}(\omega) = \sum_{j=1}^G \gamma_j e^{-2i \sum_{m=1}^j k_m L_m}, \quad (50)$$

where γ_j is the local reflection coefficient at the j -th junction. The coherent power reflection $|\Gamma_{\text{net}}(\omega)|^2$ is:

$$|\Gamma_{\text{net}}(\omega)|^2 = \sum_{j=1}^G \gamma_j^2 + 2 \sum_{1 \leq j < j' \leq G} \gamma_j \gamma_{j'} \cos(2\Delta\theta_{j,j'}) \exp(-2\Delta\kappa_{j,j'}), \quad (51)$$

where the cumulative phase shift and attenuation between junctions j and j' are $\Delta\theta_{j,j'} = \text{Re}\{\sum_{m=j+1}^{j'} k_m L_m\}$ and $\Delta\kappa_{j,j'} = \text{Im}\{\sum_{m=j+1}^{j'} k_m L_m\}$.

Physical Hypothesis. The Random-Phase Approximation (RPA) assumption, requiring length variance $\sigma_L \gg \lambda/(2\pi)$, is a modeling hypothesis justified by the stochastic variability observed in morphometric studies. While individual vascular networks are deterministic structures, the complex spatial constraints of space-filling tissue beds produce an effective randomization of segment lengths across the ensemble. Full coherent validation using measured morphometry is performed in Paper II [3]. Let us therefore model the segment lengths as independent random variables with mean $\langle L \rangle$ and variance σ_L^2 . If the length variance is large compared to the wave coherence length, i.e., $\sigma_L \gg \lambda/(2\pi)$ (where $\lambda = 2\pi/\text{Re}\{k\}$ is the wavelength), the phase factors of the cross-terms fluctuate rapidly.

Applying the Random-Phase Approximation (RPA), the ensemble expectation value of the interference terms vanishes identically:

$$\mathbb{E} [\cos(2\Delta\theta_{j,j'})] = 0 \quad \text{for } j \neq j'. \quad (52)$$

Consequently, the ensemble-averaged net power reflection is simply the sum of individual junction reflection powers:

$$\mathbb{E} [|\Gamma_{\text{net}}(\omega)|^2] = \sum_{j=1}^G \gamma_j^2. \quad (53)$$

For a symmetric tree where all junctions have identical local reflection power γ^2 , this expectation simplifies to:

$$\mathbb{E} [|\Gamma_{\text{net}}(\omega)|^2] = G\gamma^2. \quad (54)$$

B.2 Connection to the Multiplicative Geometric Penalty

Our main text metabolic cost model represents the global wave reflection penalty through the multiplicative geometric form:

$$C_{\text{wave}} = 1 - (1 - \gamma^2)^G. \quad (55)$$

Performing a Taylor expansion of this geometric penalty in powers of the local reflection coefficient γ^2 (valid for the physiologically typical weak-reflection regime where amplitude reflection is $\gamma \approx 0.05$):

$$C_{\text{wave}} = 1 - \left(1 - G\gamma^2 + \frac{G(G-1)}{2}\gamma^4 - \dots\right) = G\gamma^2 - \mathcal{O}(\gamma^4). \quad (56)$$

For typical physiological values $\gamma \approx 0.05$ (meaning $\gamma^2 \approx 0.0025$) and $G = 11$, the second-order truncation term $G(G-1)\gamma^4/2 \approx 3.4 \times 10^{-4}$ represents only a $\sim 1.25\%$ correction to the leading-order effect $G\gamma^2 \approx 0.0275$. This explicitly confirms the validity of the first-order truncation in physiological regimes.

Thus, the multiplicative geometric penalty C_{wave} is mathematically equivalent to the ensemble-averaged incoherent scattering power reflection to first order in γ^2 :

$$C_{\text{wave}} = \mathbb{E} [|\Gamma_{\text{net}}(\omega)|^2] - \mathcal{O}(\gamma^4). \quad (57)$$

Spectral averaging over a broad-band cardiac pulse containing multiple harmonics is expected to further suppress residual coherent interference. Full coherent wave simulations (Paper II [3]) confirm that the incoherent geometric penalty provides an accurate representation of mean reflection losses, with deviation $\langle |\Delta\alpha^*| \rangle < 0.01$ across physiological parameter ranges.

C Appendix C: Sensitivity Analysis of the Transition Mass M^*

The critical allometric mass M^* marking the Womersley transition (Theorem 15) is grounded in the reference state of a healthy adult human. A purely dimensional scaling relation, $M \propto \text{Wo}^4$, yields an order-of-magnitude upper bound ($M \sim 5\text{--}10\text{ g}$). However, the precise theoretical centroid $M^* \approx 0.84\text{ g}$ is determined numerically as the inflection point of the full sigmoidal wave-cost transition $\alpha^*(M)$ governed by the symmetry-motivated fluid threshold $\text{Wo}_c^{\text{fluid}} = \sqrt{3} \approx 1.732$. While we report this numerical value with precision, the physical significance of M^* resides in its order of magnitude ($\sim 1\text{ g}$) and its role as a universal topological separator.

However, Wo_0 inherits a sensitivity to physiological parameters: $\text{Wo}_0 \propto r_0 \sqrt{f_H/\nu}$. Consequently, the predicted mass threshold scales as $M^* \propto f_H^{-2} \nu^2$. To assess the robustness of the minimax allometric threshold, we compute the shift in M^* under a $\pm 10\%$ perturbation of the heart rate:

- **Tachycardic shift ($+10\% f_H$):** The transition mass shifts downward to $M^* \approx 0.69\text{ g}$.
- **Bradycardic shift ($-10\% f_H$):** The transition mass shifts upward to $M^* \approx 1.04\text{ g}$.

While the absolute threshold exhibits an asymmetric $[-17\%, +23\%]$ sensitivity to heart rate fluctuations (as expected from the inverse-square relation $M^* \propto f_H^{-2}$ where $1.1^{-2} \approx 0.83$ and $0.9^{-2} \approx 1.23$), the “logarithmic sharpness” of the jump remains invariant. Even under extreme physiological noise, the theory predicts that the metabolic transition remains confined to the sub-gram to few-gram range, insensitive to physiological variability, distinguishing it clearly from stochastic metabolic noise.

Finally, we address the robustness of the thin-wall approximation. In the terminal microvasculature ($h/r \approx 0.42$), the classical Moens-Korteweg model requires a thick-wall correction via

the Lamé elastic solution. However, as established in the Supplemental Material, the network exhibits a *Topological Shielding Principle*: because the wave-reflection penalty vanishes as $W_o \rightarrow 0$, the architectural impact of modelling errors in distal generations is asymptotically suppressed. Numerical sensitivity analysis (see Supplemental Material, Table S3) confirms that the global minimax α^* is structurally decoupled from the breakdown of the thin-wall model, ensuring the universality of the incommensurability principle across the entire vascular hierarchy.

References

- [1] Sami F. Noujaim, Elena Lucca, Viviana Muñoz, Dharmendra Persaud, Omer Berenfeld, Frits L. Meijler, and José Jalife. From mouse to whale: A universal scaling relation for the PR interval of the electrocardiogram of mammals. *Circulation*, 110(19):2802–2808, 2004. doi: [10.1161/01.CIR.0000146785.15995.67](https://doi.org/10.1161/01.CIR.0000146785.15995.67).
- [2] R. Marchesi. Beyond Murray’s law: Non-universal branching exponents from vessel-wall metabolic costs. *arXiv preprint*, 2026. doi: [10.48550/arXiv.2603.13687](https://doi.org/10.48550/arXiv.2603.13687).
- [3] R. Marchesi. A unified variational principle for branching transport networks: Wave impedance, viscous flow, and tissue metabolism. *arXiv preprint*, 2026. doi: [10.48550/arXiv.2603.14691](https://doi.org/10.48550/arXiv.2603.14691).
- [4] R. Marchesi. The dynamic origin of Kleiber’s law and the generalized metabolic scaling theorem. *arXiv preprint*, 2026. doi: [10.48550/arXiv.2604.10476](https://doi.org/10.48550/arXiv.2604.10476).
- [5] A. Taylor et al. Systematic review and meta-analysis of Murray’s Law in the coronary arterial circulation. *Am. J. Physiol. Heart Circ. Physiol.*, 327(1):H182–H190, 2024. doi: [10.1152/ajpheart.00142.2024](https://doi.org/10.1152/ajpheart.00142.2024).
- [6] L. Onsager. Reciprocal relations in irreversible processes. *Phys. Rev.*, 37:405–426, 1931. doi: [10.1103/PhysRev.37.405](https://doi.org/10.1103/PhysRev.37.405).
- [7] R. Lande and S. J. Arnold. The measurement of selection on correlated characters. *Evolution*, 37(6):1210–1226, 1983. doi: [10.2307/2408842](https://doi.org/10.2307/2408842).
- [8] Van M. Savage, James F. Gillooly, William H. Woodruff, Geoffrey B. West, Andrew P. Allen, Brian J. Enquist, and James H. Brown. The predominance of quarter-power scaling in biology. *Functional Ecology*, 18(2):257–282, 2004. doi: [10.1111/j.0269-8463.2004.00856.x](https://doi.org/10.1111/j.0269-8463.2004.00856.x).
- [9] Aharon Ben-Tal, Laurent El Ghaoui, and Arkadi Nemirovski. *Robust Optimization*. Princeton University Press, 2009. doi: [10.1515/9781400831050](https://doi.org/10.1515/9781400831050).
- [10] G. S. Kassab, C. A. Rider, N. J. Tang, and Y. C. Fung. Morphometry of pig coronary arterial trees. *Am. J. Physiol.*, 265:H350–H365, 1993. doi: [10.1152/ajpheart.1993.265.1.H350](https://doi.org/10.1152/ajpheart.1993.265.1.H350).
- [11] A. Forjaz, M. Costa, C. Oliveira, P. A. Gensbigler, L. Dequiedt, V. Queiroga, S. Joshi, W. Foster, M. Wyart, M. R. Grafe, O. J. T. McCarty, G. S. Bever, J. O. Lo, B. Menard, S. X. Sun, A. L. Kiemen, and D. Wirtz. Whole organism 3D mapping reveals universal branching topology and biophysical optimization governs vascular and nervous system development. *bioRxiv preprint*, 2026. doi: [10.64898/2026.04.10.717729](https://doi.org/10.64898/2026.04.10.717729).
- [12] Z. L. Jiang, G. S. Kassab, and Y. C. Fung. Diameter-defined Strahler system and connectivity matrix of the pulmonary arterial tree. *J. Appl. Physiol.*, 76(2):882–892, 1994. doi: [10.1152/jappl.1994.76.2.882](https://doi.org/10.1152/jappl.1994.76.2.882).

- [13] W. Huang, R. T. Yen, M. McLaurine, and Y. C. Fung. Morphometry of the human pulmonary vasculature. *J. Appl. Physiol.*, 81:2123–2137, 1996. doi: [10.1152/jappl.1996.81.5.2123](https://doi.org/10.1152/jappl.1996.81.5.2123).
- [14] S. Rossitti and J. Löfgren. Vascular dimensions of the cerebral arteries follow the principle of minimum work. *Stroke*, 24(3):371–377, 1993. doi: [10.1161/01.STR.24.3.371](https://doi.org/10.1161/01.STR.24.3.371).
- [15] T. Luo, T. J. Gast, T. J. Vermeer, and S. A. Burns. Retinal vascular branching in healthy and diabetic subjects. *Invest. Ophthalmol. Vis. Sci.*, 58(5):2685–2694, 2017. doi: [10.1167/iovs.17-21653](https://doi.org/10.1167/iovs.17-21653).
- [16] K. A. McCulloh, J. S. Sperry, and F. R. Adler. Water transport in plants obeys Murray’s law. *Nature*, 421:939–942, 2003. doi: [10.1038/nature01444](https://doi.org/10.1038/nature01444).
- [17] D. D. Quintana, S. E. Lewis, Y. Anantula, J. A. Garcia, S. N. Sarkar, J. Z. Cavendish, C. M. Brown, and J. W. Simpkins. The cerebral angiome: High resolution MicroCT imaging of the whole brain cerebrovasculature in female and male mice. *NeuroImage*, 202:116109, 2019. doi: [10.1016/j.neuroimage.2019.116109](https://doi.org/10.1016/j.neuroimage.2019.116109).
- [18] K. S. Burrowes, P. J. Hunter, and M. H. Tawhai. Anatomically based finite element models of the human pulmonary arterial and venous trees including supernumerary vessels. *J. Appl. Physiol.*, 99:731–738, 2005. doi: [10.1152/japplphysiol.01033.2004](https://doi.org/10.1152/japplphysiol.01033.2004).
- [19] C. L. Walsh, P. Tafforeau, W. L. Wagner, D. J. Jafree, A. Bellier, et al. Imaging intact human organs with local resolution of cellular structures using hierarchical phase-contrast tomography. *Nat. Methods*, 18:1532–1541, 2021. doi: [10.1038/s41592-021-01317-x](https://doi.org/10.1038/s41592-021-01317-x).
- [20] M. Chourrout, A. Keenlyside, E. Wanjau, Y. Balbastre, E. Yagis, J. Brunet, D. Stansby, K. Engel, X. Gui, J. Thönnißen, T. Dickscheid, L. Lamalle, A. Bellier, U. Vivekananda, P. Tafforeau, P. D. Lee, and C. L. Walsh. The extremely brilliant brain: An isotropic microscale human brain dataset. *bioRxiv preprint*, 2026. doi: [10.64898/2026.01.27.702076](https://doi.org/10.64898/2026.01.27.702076).
- [21] H. Uhlirova, P. Tian, K. Kılıç, M. Thunemann, V. B. Sridhar, H. Bartsch, A. M. Dale, A. Devor, and P. A. Saisan. Neurovascular network explorer 2.0: A database of 2-photon single-vessel diameter measurements from mouse SI cortex in response to optogenetic stimulation. *Front. Neuroinform.*, 11:4, 2017. doi: [10.3389/fninf.2017.00004](https://doi.org/10.3389/fninf.2017.00004).
- [22] J. L. Lucitti, K. Tobita, and B. B. Keller. Arterial hemodynamics and mechanical properties after circulatory intervention in the chick embryo. *J. Exp. Biol.*, 208:1877–1885, 2005. doi: [10.1242/jeb.01574](https://doi.org/10.1242/jeb.01574).
- [23] J. Bennett. Murray’s law as an entropy-per-information-cost extremum. *arXiv preprint*, 2025. doi: [10.48550/arXiv.2511.04022](https://doi.org/10.48550/arXiv.2511.04022).
- [24] J. Bennett. A single-index theory of optimal branching. *arXiv preprint*, 2025. doi: [10.48550/arXiv.2511.19915](https://doi.org/10.48550/arXiv.2511.19915).
- [25] I. Prigogine. *Introduction to Thermodynamics of Irreversible Processes*. Wiley, New York, 3rd edition, 1967. ISBN 0470699280. URL <https://www.worldcat.org/isbn/0470699280>.
- [26] Yunlong Huo and Ghassan S. Kassab. Pulsatile blood flow in the entire coronary arterial tree: theory and experiment. *Am. J. Physiol. Heart Circ. Physiol.*, 291:H1074–H1087, 2006. doi: [10.1152/ajpheart.00200.2006](https://doi.org/10.1152/ajpheart.00200.2006).

- [27] X. Guo and G. S. Kassab. Distribution of stress and strain in the porcine coronary arterial tree. *Am. J. Physiol. Heart Circ. Physiol.*, 284(1):H53–H62, 2003. doi: [10.1152/ajpheart.00518.2002](https://doi.org/10.1152/ajpheart.00518.2002).
- [28] Dong Hu and David Cai. Adaptation and optimization of biological transport networks. *Physical Review Letters*, 111:138701, 2013. doi: [10.1103/PhysRevLett.111.138701](https://doi.org/10.1103/PhysRevLett.111.138701).
- [29] Henrik Ronellenfitsch and Eleni Katifori. Global optimization, local adaptation, and the role of growth in distribution networks. *Physical Review Letters*, 117:138301, 2016. doi: [10.1103/PhysRevLett.117.138301](https://doi.org/10.1103/PhysRevLett.117.138301).
- [30] Robin Fåhræus and Torsten Lindqvist. The viscosity of the blood in narrow capillary tubes. *Am. J. Physiol.*, 96:562–568, 1931.